

LakeCat

covers lakecat (0.2.0-c80a309e)

An Engine-Close Catalog Foundation for QueryGraph

Alexy Khrabrov

chiefscientist.org

&

Codex with ChatGPT 5.5

Contents

1	Preface	2
2	Catalogs, Iceberg, and What LakeCat Adds	2
2.1	What a catalog is	2
2.2	What Iceberg makes the catalog responsible for	3
2.3	What LakeCat adds without changing Iceberg	4
3	The Vocabulary	4
4	The Boundary Model	7
5	Where LakeCat Stands Today	9
6	Why the Catalog Stays Thin	10
7	The Architecture	12
7.1	The CatalogStore seam	12
7.2	The read path	13
7.3	The commit path	14
7.4	The durable spine	16
8	The Siblings and the Engine Path	16
8.1	Grust: the graph boundary	16
8.2	TypeSec: the authorization boundary	17
8.3	Sail and the v3→v4 path	17
8.4	The semantic handoff: Croissant, CDIF, OSI, OpenLineage	19
8.5	The full stack	20
8.6	Implementation shape	20
9	Workflow Examples	21
9.1	Starting The Catalog	21
9.2	Registering The Warehouse Shape	22
9.3	Storage Profiles And Credential Roots	23
9.4	A PySpark User Reads Iceberg	24
9.5	A Notebook Requests Credentials	25
9.6	A View Becomes Part Of The Catalog Story	25
9.7	QueryGraph Bootstrap	26
9.8	Draining The Outbox	28
9.9	An Agentic QGLake Flow	28
10	Operating The Book's Example System	29
11	Standards And Engine Boundary Decision Record	30
11.1	Catalog Concepts In Plain Terms	30
11.2	Decision Matrix	33
11.3	The Proper-Noun Test	35
11.4	Why The Engine Owns The Truth	35
11.5	Workflow Consequences	36
12	First Release Readiness	36
13	Appendix A: Replay Admission Matrix	41
14	Appendix B: Credential And Redaction Contract	43

15	Appendix C: View Receipt-Chain Rules	45
16	Appendix D: QGLake Handoff Proof Map	47
17	Future Work	49

1 Preface

LakeCat is a Rust-native, Iceberg-compatible catalog foundation for QueryGraph. It begins from one conservative claim and one ambitious one. The conservative claim: an ordinary Iceberg REST catalog must keep working for ordinary engines — Spark, Flink, Trino, DuckDB, PyIceberg, Sail — with no new protocol to learn. The ambitious claim: the *next* catalog also has to be a governed control plane for Rust-first planning, semantic graph handoff, lineage, and agent access.

LakeCat holds both at once by drawing a sharp line. It keeps Iceberg compatibility at the boundary and moves the work that needs deep table-format knowledge to the engine that can reason about it. In this repository the catalog is **LakeCat** and the engine-side lakehouse implementation is **Sail**. Graph behavior belongs to **Grust**; governance and capability proof belong to **TypeSec**; the end-to-end integration target is **QueryGraph**.

This book builds from first principles. Chapter 2 explains what a catalog is, what Iceberg makes the catalog responsible for, and what LakeCat adds without changing the table format. Chapter 3 fixes the vocabulary — the single hardest thing in a layered system — so every later term has an owner. Chapter 4 draws the boundary model: which repository owns which concept, and which ideas might someday become neutral, standardizable profiles. The middle chapters walk the live architecture: the service spine, the read path, the commit path, and the durable store. The final chapters cover the engine boundary and the v3→v4 path, the graph/security/lineage handoffs, the QueryGraph/QGLake acceptance flow, worked examples, and the first-release scope.

Each concept is defined once and then used. Where a later chapter needs a term, it links back rather than restating it.

2 Catalogs, Iceberg, and What LakeCat Adds

2.1 What a catalog is

A data catalog is often described as a place that lists datasets. That is true but too small. A real catalog is the control plane between names, storage, metadata, identity, and intent. At minimum it answers four questions:

1. What table does this name mean?
2. Where is its current metadata?
3. Who may read, write, plan, or administer it?
4. What changed, when, and under whose authority?

In a traditional database the catalog is embedded in the engine: one system owns table definitions, statistics, permissions, and the transaction log. A lakehouse pulls that apart. Data files live in object storage, metadata files live beside them, and several engines read and write the same tables.

The catalog becomes the agreement point — it maps a logical table name to the current metadata pointer and arbitrates updates to that pointer.




Figure 1: Diagram 1

That pointer is deceptively important. If the catalog points at metadata version 17, the table is version 17. If a writer prepares version 18 and wins the compare-and-swap, the table becomes 18; if it loses, nothing partially changes. The catalog is not the table format, but it is where table history becomes visible and durable. For human analytics this sounds like bookkeeping. For agentic systems it is a trust boundary: a catalog can know the principal, warehouse, namespace, table, snapshot, requested columns, row restriction, storage profile, and policy receipt — and if it captures that before planning and commits it after state changes, it becomes a governed control plane rather than a passive address book.

2.2 What Iceberg makes the catalog responsible for

Apache Iceberg is a table format for large analytic tables. Its core idea is to put the table’s truth in explicit metadata files so engines can plan reads and validate writes without directory listing or fragile storage conventions. A current metadata file names schemas, partition specs, sort orders, snapshots, and properties; snapshots point to manifest lists; manifest lists point to manifests; manifests describe data and delete files.

The catalog’s role in Iceberg is intentionally narrow. Standard clients must be able to load table metadata, create namespaces and tables, commit changes, and sometimes receive credentials or scan tasks — over a documented REST shape:

```
GET /v1/config
GET /v1/{prefix}/namespaces
POST /v1/{prefix}/namespaces/{ns}/tables
GET /v1/{prefix}/namespaces/{ns}/tables/{table}
POST /v1/{prefix}/namespaces/{ns}/tables/{table} # commit: requirements +
updates
```

The rule that governs everything else in this book: **a normal client must never have to call a non-standard endpoint to read an ordinary table**. If it does, compatibility is already broken.

2.3 What LakeCat adds without changing Iceberg

LakeCat’s promise is *compatibility first, evidence second, semantics above the catalog*. A Spark or PyIceberg client sees an Iceberg REST catalog. A QueryGraph or governed-agent client may ask for richer proof. Both share the same table because the portable truth stays in Iceberg metadata and the extra evidence sits beside it — never inside the table format.




Figure 2: Diagram 2

Iceberg metadata stays pristine. Policy, graph, lineage, and agent state are *derived* control-plane or graph data. The table remains an Iceberg table.

3 The Vocabulary

A layered system fails the moment its words blur. The most common mistake is to call every useful LakeCat feature an “Iceberg extension,” which makes the standard boundary too large. The test is simple: **ask what breaks if a client knows nothing about LakeCat**. If a PySpark job cannot load, commit, or drop a table without the concept, it belongs to standard compatibility. If PySpark keeps working but operators, governed agents, or QueryGraph gain stronger evidence, the concept is an additive surface.

Every term in this book falls into one of six categories.




Figure 3: Diagram 3

Standard Iceberg parlance — the words Iceberg already owns: catalog, namespace, table identifier, current metadata location, snapshot, manifest list, manifest, data file, delete file, schema and partition evolution, optimistic commit, REST compatibility. LakeCat must implement these faithfully. Changing their meaning means losing compatibility.

LakeCat implementation — how this Rust catalog satisfies the contract reliably: the service spine, the CatalogStore trait, the Turso-backed durable store, normalized idempotency rows, pointer logs, audit rows, outbox rows, redaction rules, and replay validators. These make ordinary Iceberg behavior atomic, inspectable, and replayable. They are not extensions; they are good engineering behind a standard surface.

LakeCat optional extensions — additive APIs beside the standard path: management inventory, commit-history inspection, view proof, credential-root posture, replay verification, OpenLineage projection, and QueryGraph/QGLake bootstrap bundles. They help operators, agents, and QueryGraph without becoming hidden requirements for ordinary table access.

TypeSec governance — authority and receipt semantics: who may act, capability proof, TypeDID envelopes, ODRL-derived restrictions, and agent posture. TypeSec *decides*; LakeCat *carries and records* the decision.

QueryGraph semantics — composition above the catalog: Croissant, CDIF, OSI, ODRL, OpenLineage, and Grust graph projections built from the governed source of truth.

Sail engine truth — table-format interpretation, metadata-as-data, scan planning, commit validation, and typed v4 behavior. The catalog binds its proof to engine truth rather than reimplementing the format.

The single reference below replaces the per-claim restatements that earlier drafts repeated; read each LakeCat claim through its category.

Claim	Standard reading	Category	Portable idea (if any)
Rust service / catalog spine	Iceberg needs a catalog authority, not a language	LakeCat implementation	“A catalog can prove what it committed, planned, vended, and emitted”
Turso-backed durable store	Iceberg needs durable state + atomic pointer movement, not a named DB	LakeCat implementation	CAS, exact-retry, row/content binding, pointer-history proof
REST namespace/table routes + commit CAS	Standard compatibility	Standard Iceberg	— (must stay ordinary)
Idempotency, pointer logs, audit/outbox, replay validation	Mostly outside the table contract	LakeCat implementation / extension	Stable catalog-event identity; scoped replay admission
Governed scan with policy receipt	Iceberg gives scan inputs, not receipts	TypeSec governance	Engine-proved effective projection + predicate
Credential vending posture	Catalog-adjacent, not ordinary table semantics	TypeSec governance	Bounded, redacted, engine-neutral credential profile
QueryGraph / QGLake / OpenLineage handoff	Not required for table access	QueryGraph semantics	Redacted, replayable lineage/view/commit evidence
Typed Iceberg v4 behavior	Belongs to the format as it evolves	Sail engine truth	Engine-owned v4 interpretation, not catalog JSON parsing

4 The Boundary Model

The vocabulary tells you which category a word lives in. The boundary model says which *repository* owns the work and why. LakeCat is deliberately thin: it keeps identity, tenancy, metadata-pointer state, policy gates, idempotent commits, and integration events — and delegates everything reusable.




Figure 4: Diagram 4

The ownership rule, stated once:

Concern	Owner	LakeCat keeps only
Iceberg format, manifests, scan planning, pruning, delete handling, v4	Sail	The call into Sail and the proof binding its result
Graph schema, taxonomy, traversal, stores, Cypher	Grust	The catalog-facing sink/projection boundary
Authorization, policy composition, capabilities, TypeDID, credential decisions	TypeSec	The request for a decision and the persisted receipt
Croissant/CDIF/OSI/ODRL/OpenLineage composition, agent workflows	QueryGraph	The governed bootstrap bundle it emits
Identity, tenancy, pointer CAS, idempotency, audit, outbox, replay	LakeCat	All of it — this is the thin catalog

This is also the answer to the recurring “is it an extension or a standard?” question. LakeCat is opinionated in code and modest in standardization. “Use Rust,” “use Turso,” “use TypeSec,” “import QGLake,” “project into Grust” are project choices, not proposals. The portable ideas are narrower and stated without product names: *reject idempotency drift; record redacted pointer history; emit transactional catalog-event identity; admit only scoped replay evidence; prove a governed scan was narrowed by an engine*. A proposal that forced every Iceberg catalog to understand QueryGraph, TypeSec, Grust, or Turso would narrow the ecosystem; a proposal that says “a catalog may publish redacted, replayable commit evidence” leaves room for many implementations. Prove the stronger shape locally, then extract only the database-neutral, policy-neutral, engine-neutral part.

5 Where LakeCat Stands Today

LakeCat is not only a design. The implementation already has a Rust service/catalog spine, a Turso-backed durable store, Iceberg REST-compatible namespace and table paths, hardened commit and replay evidence, governed scan and credential proof, and QueryGraph/QGLake handoff surfaces. Read each of those through its category (Chapter 3) rather than flattening them into “Iceberg extensions”: some are ordinary catalog implementation, some are optional LakeCat APIs, some are TypeSec governance proof, and only a few are plausible future standardization candidates.

The most important word is **beside**. New capability sits beside the Iceberg REST path, never in front of it. A standard client never has to present a TypeDID, parse a QueryGraph bundle, read a Grust edge, or inspect an OpenLineage receipt to load a normal table. So the same catalog shows a different face to each caller:

Caller	What they see
PySpark / PyIceberg user	Ordinary Iceberg REST. Configure a catalog, create a namespace, write and load tables. Pointer-log, audit, and outbox rows exist but are invisible to table semantics.
Platform operator	A hardened catalog transaction log: idempotency outcomes, pointer movement, redacted conflict proof, storage-profile and credential posture, pending outbox, replay-validation failures.
Governed agent	A narrowed access path: TypeSec decides the capability, LakeCat binds the receipt to the catalog action, Sail plans the effective scan, the agent gets bounded work instead of broad credentials.
QueryGraph importer	A proof-bearing bootstrap: table, view, management, credential, scan, commit-history, OpenLineage, and graph-import anchors that must line up before the semantic layer is accepted.

That separation also keeps standardization honest. The interesting portable shapes are never “LakeCat uses Turso” or “QueryGraph imports a bundle”; they are smaller behaviors — idempotent commit replay, catalog pointer history, governed credential vending, proof-carrying scan planning, redacted conflict evidence, lineage receipt binding. The posture, in five rules:

1. Keep base Iceberg REST behavior strict and boring.
2. Keep LakeCat proof surfaces optional for the clients that need them.
3. Keep TypeSec and QueryGraph semantics out of Iceberg metadata.
4. Use Sail for reusable table-format and planning semantics.
5. Promote only small, interoperable proof profiles once real workflows prove multiple engines and catalogs need them.

6 Why the Catalog Stays Thin

The most dangerous failure mode for a “smart” catalog is becoming a partial engine. It starts innocently — validate a schema here, expand a manifest list there, peek at a format-version field, check a delete file. Each small parser looks cheaper than an engine call. Over time the catalog grows a second Iceberg implementation with weaker tests, fewer real execution users, and quiet drift from how the planner actually behaves.




Figure 5: Diagram 5

LakeCat takes the opposite path: **the catalog owns the transaction; Sail owns the table semantics**. Sail is the right home because it is Rust-native and already holds the structures this needs — generated Iceberg REST models, table providers, manifest pruning, metadata-as-data paths, commit plumbing, and format-version handling. Anything that needs field-id binding, schema/partition evolution, manifest metrics, delete association, row lineage, v4 metadata trees, or snapshot/branch selection moves toward Sail.

This matters because intermediation, left passive, *loses* information. A pass-through catalog sees only the table name, the pointer, the caller, and maybe a credential request; the engine sees

schema, manifests, statistics, deletes, and filters; governance sees policy; lineage sees an after-the-fact event. Each system gets a shard of the truth, and the operational cost is concrete:

- policy can be checked before access but not carried *into* scan planning;
- credentials can be vended without proving why a raw-credential exception was allowed;
- lineage can say something happened but cannot bind to the exact governed plan, snapshot, policy, and metadata.

The next generation of callers — agents, notebooks, services, model pipelines — needs stronger guarantees that live precisely between catalog state and engine planning: a policy enforced before a scan is planned; column restrictions that narrow the projection before file tasks exist; policy-derived row predicates that become mandatory filters; a stateless `fetchScanTasks` that cannot widen a prior governed plan; short-lived, scoped, audited credentials; idempotent commit retries; and graph/lineage side effects that reflect committed state rather than best-effort handler work. A catalog too far from the engine can only check a policy and hand back a pointer, hoping the client preserves the restriction. LakeCat keeps the catalog boundary thin and binds its proof to Sail’s plan so the restriction is real.

7 The Architecture

Thin does not mean trivial. LakeCat owns the durable catalog state that must be correct even when external sinks are down, and nothing more. Its responsibilities are exactly:

- serve the Iceberg REST Catalog API for standard clients;
- model projects, warehouses, namespaces, tables, views, and storage profiles;
- persist metadata pointers and compare-and-swap commit history;
- validate idempotency keys and replay only matching commit bodies;
- resolve request identity from headers, bearer tokens, agents, and TypeDID envelopes;
- ask TypeSec for authorization decisions and persist the receipts;
- route scan planning and commit preparation through Sail;
- record audit and outbox events inside the catalog transaction;
- drain committed events to Grust and OpenLineage sinks;
- publish a QueryGraph bootstrap bundle.

Everything else is deliberately excluded — LakeCat does not invent a table format, fork manifest pruning, own graph traversal, decide security semantics, or author QueryGraph’s business model. Those belong to Sail, Grust, TypeSec, and QueryGraph respectively (Chapter 4).

7.1 The CatalogStore seam

All durable state sits behind one trait, so the memory store (for tests and embedded use) and the Turso store (for durable local deployments) are interchangeable. The interesting methods are the table lifecycle and the idempotent replay hook:

```
#[async_trait]
pub trait CatalogStore: Send + Sync + 'static {
    async fn create_namespace(&self, warehouse: &WarehouseName, namespace:
Namespace)
```

```

        -> LakeCatResult<()>;
    async fn list_tables(&self, warehouse: &WarehouseName) ->
LakeCatResult<Vec<TableRecord>>;
    async fn create_table(&self, table: TableRecord) ->
LakeCatResult<TableRecord>;
    async fn load_table(&self, ident: &TableIdent) -> LakeCatResult<TableRecord>;

    // Optimistic pointer advancement: the catalog transaction.
    async fn commit_table(&self, ident: &TableIdent, commit: TableCommit)
        -> LakeCatResult<TableRecord>;

    // Exact idempotent replay: same key + same request hash -> stored response.
    async fn replay_table_commit(
        &self,
        ident: &TableIdent,
        idempotency_key: &str,
        idempotency_request_hash: &str,
    ) -> LakeCatResult<Option<TableRecord>>;

    // Compact pointer-log history for operators and QueryGraph.
    async fn table_commit_records(&self, ident: &TableIdent, start: u64, end:
Option<u64>)
        -> LakeCatResult<Vec<TableCommitRecord>>;
}

```

7.2 The read path

A read begins like a standard catalog request and ends with a *governed* Sail plan. The policy restriction becomes part of planning, not a note beside it.




Figure 6: Diagram 6

The governing rule is **narrow, never widen**. An empty client projection under a column restriction means *the allowed columns*; a client projection may narrow further but cannot widen. LakeCat records both the requested and the effective projection (and stats fields) as replay evidence, and recomputes the restriction on every stateless `fetchScanTasks` so a stale token cannot expand back to all columns. Outbox admission enforces that the governed scan replay carries the same read-restriction at the top level and inside `authorization-receipt.context`, a nonblank purpose, and a positive `max-credential-ttl-seconds`, and rejects unknown fields — so graph, lineage, and QGLake evidence can never inherit a claim the receipt did not capture.

7.3 The commit path

The write path follows the same principle: LakeCat owns the transaction, Sail owns Iceberg validation and metadata assembly.




Figure 7: Diagram 7

Three invariant groups make this safe, each enforced once:

- **Metadata-write safety.** The new object must be a concrete child of the table's matched storage-profile prefix — never the current pointer, never the profile root, never a path with literal or percent-encoded dot segments, query strings, fragments, userinfo, or credential markers. Writes are create-only; if the store CAS loses after a write, LakeCat makes a bounded retry to delete the orphaned object and otherwise preserves the original conflict.
- **Idempotency.** Keys are 1–128 ASCII chars (Idempotency-Key / x-lakecat-idempotency-key, matching if both present). The same key + same body replays the stored response before any Sail call or write; the same key + a different body conflicts. Only audit-safe hashes are stored — never raw keys or secrets. Ordinary clients may commit without a key; they simply get no replay proof.
- **Redaction.** Every commit error — pointer conflict, overwrite, prefix mismatch, cleanup failure, backend error — reports sha256: hashes of the metadata location and failure detail, never raw object paths, storage roots, profile ids, or backend text. A conflict still looks like an ordinary Iceberg conflict to the client.

Commit records carry compact evidence — Iceberg format version, current snapshot id, request/response/idempotency/policy hashes, principal — so operators and QueryGraph can answer audit questions from the pointer log without parsing full metadata. Operators read it through a governed endpoint:


```
curl -s -H 'x-lakecat-principal: operator@example.com' \
  http://127.0.0.1:3000/management/v1/warehouses/local/namespaces/default/tables/
events/commits
```

That read itself enters the outbox as `table.commits-listed` and drains as lineage plus Commit graph anchors keyed by table and sequence — an auditable trail without giving anyone direct database access or turning LakeCat into a graph query engine.

7.4 The durable spine

The durable local store uses the Rust `turso` crate behind the `turso-local` feature. Object storage remains the source of Iceberg metadata files; Turso holds the atomic pointer and the control-plane memory: projects and warehouses, storage profiles, namespaces and tables, the metadata pointer log, idempotency records, soft deletes, policy bindings, audit events, and outbox events. This mirrors the Iceberg contract — metadata files describe the table; catalog state decides which file is current.

Two disciplines keep the spine trustworthy:

- **Decoded-row binding.** Every read reconciles the decoded JSON against the row’s own key columns before returning it — a warehouse against its project and storage root, a namespace against its path, a policy binding against its table and enforced flag, a storage profile against its prefix/provider/mode, a table body against its identity. A corrupted row index can therefore never point a valid-looking record at the wrong tenant, path, or policy.
- **Strict outbox draining.** LakeCat projects a batch to graph and lineage *first*, then acknowledges the whole batch; if projection fails, nothing is acknowledged, and an under-count acknowledgement is a hard error rather than a quiet partial success. The drain refuses unknown event types (they stay pending rather than vanish), rejects event-id/payload-hash drift, and validates governed-read and commit evidence — policy-hash arrays must be non-empty full digests, top-level and receipt read-restriction must match, purpose and TTL must be present — before any sink sees the event. Malformed source evidence stays available for retry instead of being promoted into a QGLake handoff.

8 The Siblings and the Engine Path

Chapter 4 named the four sibling repositories. This chapter shows how LakeCat talks to each — what it hands off, and what it deliberately refuses to own.

8.1 Grust: the graph boundary

Catalog events naturally form a graph: a server contains projects, a project contains warehouses, a warehouse contains namespaces and credential-rooted storage profiles, a namespace contains tables, and a table has columns, snapshots, commits, policies, scan plans, principals, and lineage runs. `QueryGraph` wants that graph — but LakeCat must not become a graph database.

LakeCat emits a bounded envelope of stable catalog facts through a catalog-facing sink; Grust owns the taxonomy, stores, traversal, and Cypher:

Server CONTAINS Project	Warehouse HAS_STORAGE_PROFILE StorageProfile
Project CONTAINS Warehouse	Table GOVERNED_BY Policy
Warehouse CONTAINS Namespace	Principal CAN_PLAN ScanPlan
Namespace CONTAINS Table	Commit DERIVED_FROM Snapshot
Table HAS_COLUMN Column	LineageRun USED_BY QueryGraphModel

High-cardinality file and manifest facts stay queryable through Sail’s metadata-as-data rather than being smuggled into the event sink. Storage-profile and credential-vend events replay with redacted evidence only — `secret-ref-present`, the provider, never the secret URI — so Query-Graph can see *that* a principal attempted credential-root access without seeing the credential. The `grust-turso-local` feature proves the boundary end to end: LakeCat writes catalog events into a Grust-owned Turso graph store and Grust Cypher reads them back, with LakeCat never parsing Cypher or executing traversal.

8.2 TypeSec: the authorization boundary

LakeCat is a policy *enforcement* point, not the author of security semantics. Every externally meaningful action — config reads; namespace and table lifecycle; scan planning; commit; credential vending; policy management; graph and lineage reads — passes through TypeSec. LakeCat gathers context (principal and agent DID, bearer subject, warehouse/namespace/table, columns, snapshot, requested credential duration, purpose, active bindings), asks TypeSec for a decision, persists the receipt with audit-safe hashes, and applies the restriction *before* Sail plans.

This is where ODRL becomes operational. A policy may say a principal reads only certain columns, only for a purpose, or only under a row predicate. LakeCat parses the minimal enforceable subset — accepting camel-case, kebab-case, and JSON-LD operand forms — and **fails closed**: missing or deny-shaped operators, blank allowed-column lists, blank purposes, or disagreeing purpose sources are rejected rather than guessed. Composition and reasoning stay in TypeSec; LakeCat does not grow a parallel security language.

Credential vending is the audited exception, not the default. Governed Sail-planned reads are the norm. When credentials must issue, TypeSec checks the `credentials.issue` capability for the exact secret reference and LakeCat returns only scoped, short-lived configuration — capping TTL to the tightest `max-credential-ttl-seconds` across all policy locations, replacing any issuer-supplied `lakecat.*` evidence with catalog-derived values, and recording only hashed prefixes in audit. The replay-admission rules from Chapter 5 (matching top-level and receipt restrictions, full digests, closed field sets) apply identically here, so a credential or raw-exception event cannot drift before it reaches graph, lineage, or QGLake proof.

8.3 Sail and the v3→v4 path

Sail is a Rust-native engine — Arrow, DataFusion, generated Iceberg REST models, catalog-provider seams, manifest pruning, metadata-as-data — so LakeCat can *ask* for typed Iceberg behavior instead of parsing just enough JSON to survive. The questions LakeCat sends to Sail are the ones that require table-format knowledge:

- which field ids satisfy this projection?
- which filters are enforceable at planning time?

- which manifests and files survive pruning?
- which delete files must accompany the selected data files?
- which scan tasks are children of a governed parent plan?
- which v4 fields are known, preserved as passthrough, or not yet safe to read?

LakeCat evolves under three rules: conform to Iceberg v3 for ordinary clients; preserve unknown/emerging v4 metadata without claiming settled semantics; prefer typed Sail support the moment it exists, using JSON passthrough only as a bridge. Today the v4 bridge is deliberately narrow — when LakeCat sees `format-version: 4`, `lakecat-sail` extracts the stable envelope (table UUID, location, schema id, snapshot id, sequence number, manifest-list path, default spec, field names), can plan a governed manifest-list scan task from it, and re-validates the signed plan task on `fetchScanTasks` so a stateless fetch can neither drift to a different manifest list nor widen the governed projection. Pruning and typed metadata-tree semantics wait for Sail-owned v4 support.

The handoff between LakeCat and Sail should therefore be compact and typed:

Work item	LakeCat responsibility	Sail responsibility	Durable proof Lake-Cat should keep
Table load	Resolve warehouse, namespace, table, principal, and current metadata pointer.	Interpret table metadata and expose table status using reusable Iceberg types.	Table identity, pointer hash, format version, principal, and receipt/action context.
Governed scan	Ask TypeSec for the effective restriction and bind it to the request.	Bind projection and predicates to field ids, prune manifests and files, attach deletes, and produce scan tasks.	Requested and effective projection hashes, predicate hash, snapshot, task count, delete posture, plan hash, and receipt hash.
Fetch scan task	Check that a stateless task fetch belongs to the prior governed plan.	Reapply task interpretation without widening the plan.	Prior plan hash, task token hash, effective stats-field evidence, and receipt hash.
Commit	Own CAS, idempotency, pointer-log, audit, and outbox state.	Validate table metadata, commit requirements, format-version behavior, and future v4 typed semantics.	Expected and new pointer hashes, format version, snapshot evidence, idempotency hash, audit id, and outbox id.
Metadata-as-data	Authorize the request and record who asked.	Expose snapshots, manifests, files, deletes, partitions, and history as engine-shaped metadata views.	Metadata view name, pointer hash, result-shape hash, principal, and receipt hash.

8.4 The semantic handoff: Croissant, CDIF, OSI, OpenLineage

QueryGraph needs a semantic picture, not just physical access. LakeCat publishes one without pretending to be QueryGraph. The bootstrap bundle carries Semantic Croissant and CDIF projections (dataset/field discovery), an OSI handoff of stable anchors, ODRL artifacts and TypeSec policy context, OpenLineage events for catalog changes and plans, a Grust-ready graph envelope, and a manifest that hashes every artifact. The boundary is careful: LakeCat publishes stable anchors and governed source metadata; QueryGraph authors the business model. Because OpenLineage events drain from the durable outbox in `created_at,event_id` order *after* the catalog transaction, lineage reflects committed state rather than a handler’s best-effort side effect.

8.5 The full stack

When LakeCat is done, a standard engine still loads and commits tables without knowing QueryGraph exists, while governed callers get the richer path:




Figure 8: Diagram 8

8.6 Implementation shape

The workspace expresses the architecture directly:

```
crates/
lakecat-core      stable IDs, errors, time, config, content hashes
lakecat-api       Iceberg REST request/response adapters
lakecat-store     catalog state traits + Turso-backed implementation
lakecat-sail      Sail provider bridge and privileged planning client
lakecat-graph     catalog-facing Grust sink/adapters
lakecat-security  TypeSec integration and authorization receipts
lakecat-lineage   OpenLineage projection and event receipts
lakecat-querygraph Croissant/CDIF/OSI/ODRL/OpenLineage bootstrap projection
lakecat-service   axum service, middleware, auth, routing
lakecat-cli       admin, local demo, conformance, bootstrap export
```

Feature gates keep integrations honest, and embedded defaults stay safe for tests so a memory-store test never accidentally depends on a sibling repo:

```
sail-local        local Sail APIs for planning and provider integration
typesec-local     local TypeSec APIs for governance and TypeDID verification
```

grust-local	local Grust APIs for catalog graph projection
grust-turso-local	Grust's Turso backend for durable catalog graph projection
turso-local	the Turso-backed durable store

The runtime honors the same line: without `sail-local`, the deferred Sail seam *rejects* scan planning rather than fabricating an empty plan, so any real read reflects the engine that interprets Iceberg metadata, never a catalog-shaped placeholder. Standard compatibility lives at `/catalog/v1`; management APIs, `/querygraph/v1/bootstrap`, and feature-gated Sail planning sit beside it, never in front of it. (The first-release gate and dependency contract that hold this together are covered in the release chapter.)

9 Workflow Examples

The catalog is easiest to understand by watching it participate in ordinary work. LakeCat should not ask users to think about graph, lineage, security, and Sail every time they read a table. Those systems should appear when they matter: at the boundary where a name is resolved, a policy is enforced, a plan is created, credentials are withheld or issued, and a durable event is replayed.

The examples below use one table, `local.default.events`, but the pattern is the same for larger warehouses. The important point is not the exact sample data. It is the catalog role in each workflow. The dense replay and handoff checks are in the appendices so the examples can stay human-readable.

9.1 Starting The Catalog

A local operator starts LakeCat as an Iceberg REST catalog plus management surface:

```
cargo run -p lakecat-service --features sail-local,turso-local,typesec-local,grust-local
```

The standard catalog path is still `/catalog/v1`. The management and QueryGraph surfaces sit beside it:

```
/catalog/v1
/management/v1
/querygraph/v1/bootstrap
```

A simple configuration read shows the split. Standard engines care about the Iceberg endpoints. Operators and QueryGraph care about management and bootstrap surfaces.

```
curl -s http://127.0.0.1:3000/catalog/v1/config
```

The defaults intentionally separate compatibility from future capability:

```
{
  "defaults": [
    {"key": "lakecat.compatibility", "value": "iceberg-rest"},
    {"key": "lakecat.format.baseline", "value": "iceberg-v1-v3"},
    {"key": "lakecat.format.v4", "value": "extension-ready"},
    {"key": "lakecat.format.v4.bridge", "value": "json-passthrough"},
    {"key": "lakecat.format.v4.typed-sail", "value": "unavailable"}
  ]
}
```

```
  ]  
}
```

That means LakeCat can preserve and replay emerging v4 metadata through the Sail JSON bridge, but it is not claiming typed Sail v4 semantics yet. The same posture is captured in replay evidence: config defaults and overrides stay ordinary duplicate-free key/value entries, advertised endpoints prove that the standard REST and additive governed surfaces were present, and QueryGraph sees compact config proof only after service replay accepts it.

9.2 Registering The Warehouse Shape

An operator usually starts with management objects. A server groups projects. A project groups warehouses. A warehouse owns namespaces, tables, views, storage profiles, policy bindings, and the metadata pointer state that standard engines see through Iceberg REST.

```
curl -s -X PUT http://127.0.0.1:3000/management/v1/servers/prod \  
-H 'content-type: application/json' \  
-d '{  
  "display-name": "Production LakeCat",  
  "endpoint-url": "https://lakecat.example.com",  
  "properties": {"owner": "platform"}  
}'
```

```
curl -s -X PUT http://127.0.0.1:3000/management/v1/projects/resilience \  
-H 'content-type: application/json' \  
-d '{  
  "display-name": "Resilience Desk",  
  "server-id": "prod",  
  "properties": {"environment": "demo"}  
}'
```

```
curl -s -X PUT http://127.0.0.1:3000/management/v1/projects/resilience/  
warehouses/local \  
-H 'content-type: application/json' \  
-d '{  
  "display-name": "Local QGLake Warehouse",  
  "storage-root": "file:///tmp/lakecat/qglake",  
  "properties": {"querygraph": "enabled"}  
}'
```

These writes are not Iceberg table metadata. They are catalog control-plane state. LakeCat persists them durably, records authorization receipts, and writes outbox events. When the outbox drains, server, project, warehouse, storage-profile, and policy-binding changes become catalog graph events and OpenLineage receipts. QueryGraph can later learn the management shape without requiring every Iceberg client to understand it.

The replay rule is simple: management evidence is closed, attributable, scoped, and redacted. Tenant-root records bind decoded IDs back to their route scope, policy-binding replay carries the ODRL content hash, namespace/table/view lifecycle replay carries the matching action receipt,

and raw endpoint or storage-root material is replaced by full hash evidence before graph, OpenLineage, QGLake, or QueryGraph sees it.

9.3 Storage Profiles And Credential Roots

Storage profiles bind a warehouse to physical storage roots and credential issuance policy. A local profile can return scoped local file configuration. A remote profile should usually reference a secret store and require TypeSec to authorize issuance before any resolver sees the secret reference.

```
curl -s -X PUT \
  http://127.0.0.1:3000/management/v1/warehouses/local/storage-profiles/local-
events \
  -H 'content-type: application/json' \
  -d '{
    "location-prefix": "file:///tmp/lakecat/qglake/events",
    "provider": "file",
    "issuance-mode": "local-file-no-secret",
    "public-config": {"lakecat.purpose": "developer-loop"}
  }'

curl -s -X PUT \
  http://127.0.0.1:3000/management/v1/warehouses/local/storage-profiles/s3-events
\
  -H 'content-type: application/json' \
  -d '{
    "location-prefix": "s3://lakecat/events",
    "provider": "s3",
    "issuance-mode": "short-lived-secret-ref",
    "secret-ref": "vault://kv/lakecat/events",
    "public-config": {
      "lakecat.region": "us-west-2",
      "lakecat.purpose": "production-events"
    }
  }'
```

The store accepts only plainly addressed roots: no query strings, fragments, URI userinfo, literal or percent-encoded dot segments, provider/location mismatches, ambiguous longest-prefix ties, secret-looking public-config, or decorated secret-ref values. Operator-facing errors and replay evidence use hashes such as storage-profile-prefix-hash, location-prefix-hash, secret-ref-hash, and public-config-key-hash rather than echoing submitted paths or secret-like strings.

The catalog row stores the public profile and secret reference, not raw cloud keys. A later credential request is checked against TypeSec and against the effective read restriction for the target table. Agents with fine-grained table restrictions are steered to governed Sail-planned reads instead of raw credentials. Trusted humans can receive audited standard credentials only when policy allows the exception.

For production-shaped secret managers, LakeCat keeps a provider-dispatch seam instead of embedding credentials in catalog state. `vault://` can use the built-in Vault HTTP backend. `aws-sm://`, `gcp-sm://`, and `azure-kv://` can dispatch to configured provider backends after TypeSec authorizes the exact secret-ref resource, and local file-backed roots let the same dispatch shape run in a single-node loop. Returned credentials must stay inside the selected storage profile prefix, carry the policy-derived TTL cap, and preserve LakeCat-owned profile/provider/mode/principal evidence. Replay records redacted prefix and issuer hashes, not credential material.

9.4 A PySpark User Reads Iceberg

A PySpark user should not need to know about QueryGraph. They configure Spark's Iceberg REST catalog and point it at LakeCat:

```
from pyspark.sql import SparkSession

spark = (
    SparkSession.builder
        .appName("lakecat-events")
        .config("spark.sql.catalog.lakecat", "org.apache.iceberg.spark.SparkCatalog")
        .config("spark.sql.catalog.lakecat.type", "rest")
        .config("spark.sql.catalog.lakecat.uri", "http://127.0.0.1:3000/catalog/v1")
        .config("spark.sql.defaultCatalog", "lakecat")
        .getOrCreate()
)

events = spark.table("default.events")
events.select("event_id", "severity").where("severity = 'critical']").show()
```

For an unrestricted principal, the flow looks like a normal Iceberg read:

1. Spark asks LakeCat to resolve `default.events`.
2. LakeCat checks the principal and table capability.
3. LakeCat returns an Iceberg-compatible table response.
4. Spark plans the read using Iceberg metadata.

For a governed principal, the more interesting path is server-side planning. The user request still looks ordinary, but LakeCat derives the mandatory restriction before Sail sees the plan:

```
client asks:    event_id, severity, raw_payload
policy allows: event_id, severity
Sail receives: event_id, severity
```

The catalog does not trust the client to remember that. The restriction is re-applied when scan tasks are fetched, and the audit payload records the policy hash, narrowed columns, row predicate, storage location, metadata location, principal, requested stats fields, and effective stats fields. QGLake handoff proof compares planned and fetched evidence so a saved artifact cannot widen the server-derived purpose, TTL cap, allowed columns, row predicate, stats fields, or policy hashes.

9.5 A Notebook Requests Credentials

Credential vending is deliberately different from scan planning. Returning storage credentials gives the client broader power than returning a governed task list, so LakeCat treats it as an exception path.

```
curl -s \
  -H 'x-lakecat-principal: agent:triage' \
  http://127.0.0.1:3000/catalog/v1/local/namespaces/default/tables/events/
credentials
```

For an agent bound by a fine-grained restriction, LakeCat should fail closed:

```
{
  "credentials": [],
  "lakecat:credential-block-reason": "fine-grained read restriction requires
Sail-planned reads"
}
```

That empty credential response is not a missing feature. It is the intended agentic posture. The agent should ask LakeCat to plan the read through Sail, not receive raw storage reach. For a trusted human principal, policy can allow an audited raw credential exception, but the audit and replay evidence still carry the principal, table, decision, reason, TTL cap, storage-profile anchor, and redacted returned-prefix hashes.

9.6 A View Becomes Part Of The Catalog Story

Views are catalog objects too. LakeCat stores durable view records with SQL, dialect, schema version, typed columns, properties, creator, and warehouse scope. They can be managed through the management API or through warehouse-prefixed catalog routes.

```
curl -s -X PUT \
  http://127.0.0.1:3000/management/v1/warehouses/local/namespaces/default/views/
events_view \
  -H 'content-type: application/json' \
  -d '{
    "sql": "select event_id, severity from default.events where severity =
'\''critical'\''",
    "dialect": "spark-sql",
    "schema-version": 1,
    "columns": [
      {"name": "event_id", "data-type": {"type": "long"}, "nullable": false},
      {"name": "severity", "data-type": {"type": "string"}, "nullable": false}
    ],
    "properties": {"owner": "resilience-desk"}
  }'
```

This is not Iceberg business metadata glued onto a table. It is catalog state about a view object. LakeCat records view.listed, view.upserted, view.loaded, and view.dropped audit events. Outbox replay projects listing reads to namespace-scoped graph and OpenLineage evidence, and projects single-view changes and reads to catalog-facing View graph events plus LakeCat Open-

Lineage view dataset receipts. QueryGraph bootstrap can then include views with OSI hashes, store-assigned view versions, view-aware graph edges, and OpenLineage view counts.

A caller that wants optimistic view behavior includes the version it believes is current:

```
curl -s -X PUT \
  http://127.0.0.1:3000/management/v1/warehouses/local/namespaces/default/views/
events_view \
  -H 'content-type: application/json' \
  -d '{
    "sql": "select event_id from default.events where severity =
'\''critical'\''",
    "dialect": "spark-sql",
    "schema-version": 2,
    "expected-view-version": 1
  }'
```

If another writer has already advanced the view, LakeCat returns a conflict before it replaces the current view or appends a receipt. Drop can use the same guard:

```
curl -s -X DELETE \
  'http://127.0.0.1:3000/management/v1/warehouses/local/namespaces/default/views/
events_view?expected-view-version=2'
```

LakeCat also writes a compact view-version receipt in the durable store. The receipt records the stable view id, assigned version, previous version, previous receipt hash, content hash, principal, operation, and timestamp. A create, drop, and recreate sequence therefore looks like version 1 upsert, version 1 drop tombstone, version 2 upsert linked to the tombstone receipt, not two unrelated version-1 chains for the same stable view id.

QueryGraph and operators can read the compact receipt chain directly from the governed management surface:

```
curl -s \
  http://127.0.0.1:3000/management/v1/warehouses/local/namespaces/default/views/
events_view/version-receipts \
  -H 'x-lakecat-agent-did: did:example:resilience-agent'
```

The response is catalog evidence, not Iceberg table metadata. It lets QueryGraph verify the version chain, including tombstones after the current view row is gone, while keeping the richer view history model available for a future Sail-owned implementation. Appendix C captures the exact receipt-chain rules.

9.7 QueryGraph Bootstrap

QueryGraph should import LakeCat facts through a verified handoff, not by scraping service internals. The bootstrap endpoint publishes a bundle with catalog tables, views, policy bindings, graph artifacts, OpenLineage artifacts, Croissant/CDIF/OSI/ODRL projections, and a manifest that hashes what was emitted:

```
curl -s \
-H 'x-lakecat-principal: agent:querygraph-importer' \
http://127.0.0.1:3000/querygraph/v1/bootstrap \
-o target/qglake/lakecat-bootstrap.json
```

The exported graph includes the tenant spine:

```
Catalog HAS_SERVER Server
Server HAS_PROJECT Project
Project HAS_WAREHOUSE Warehouse
Warehouse HAS_NAMESPACE Namespace
Namespace HAS_TABLE Table
```

When LakeCat has durable management rows, those graph nodes come from the stored `ServerRecord`, `ProjectRecord`, and `WarehouseRecord`. Replay evidence redacts storage roots and server endpoint URLs into hashes, so `QueryGraph` can prove which management state was observed without inheriting local paths, bucket roots, query tokens, URI fragments, or userinfo. The tenant anchors and warehouse-to-namespace edges are manifest-covered, so an importer can reject a bundle whose namespace is detached from the warehouse or rebound to a different durable tenant chain.

The `QueryGraph` side verifies the bundle before importing it:

```
cd /Users/alexysrc/querygraph/qg-rust

cargo run -- lakecat-verify \
  --bundle /Users/alexysrc/lakecat/target/qglake/lakecat-bootstrap.json

cargo run -- lakecat-import \
  --bundle /Users/alexysrc/lakecat/target/qglake/lakecat-bootstrap.json \
  --output .querygraph/lakecat/import-plan.json
```

The importer checks the outer bundle hash, manifest hashes, `QueryGraph-import` compatibility hash, graph hash, `OpenLineage` hash, view receipt evidence, and the accepted receipt-chain hash for each exported view. Graph validation belongs on the `QueryGraph/Grust` side; LakeCat is responsible for producing the clean catalog-facing graph projection and the replayable anchors.

For the full local handoff, LakeCat carries a script that runs both sides without writing generated artifacts into the `QueryGraph` checkout:

```
scripts/qglake-handoff-local.sh
```

The script starts LakeCat on 127.0.0.1:18181, uses a Turso-backed local store, generates the paired QGLake bootstrap bundle and lineage-drain response, runs `qglake-verify-replay`, then runs `QueryGraph`'s `lakecat-verify` and `lakecat-import` against the same bundle. The compact `handoff-summary.json` records the catalog URL, principal, table scope, LakeCat replay status, `QueryGraph-verified` table/view counts, and semantic bundle/graph/`OpenLineage/import` hashes plus accepted standards. Appendix D is the detailed handoff proof map.

9.8 Draining The Outbox

LakeCat records side effects as durable outbox events. Draining the outbox is what turns committed catalog facts into graph and lineage receipts:

```
curl -s -X POST \
  -H 'x-lakecat-principal: agent:lineage-drainer' \
  http://127.0.0.1:3000/management/v1/lineage/drain \
  -o target/qglake/lineage-drain.json
```

A useful drain response includes delivered event types, graph projection counts, lineage projection counts, receipt hashes, and the authorization proof for the drain request itself. Reading the replay stream is privileged, so LakeCat records that the drainer was allowed to read lineage evidence.

The end-to-end result is a chain:

```
catalog write
-> audit event
-> outbox event
-> graph projection
-> OpenLineage projection
-> QueryGraph import evidence
```

If graph or lineage sinks are down, catalog state should not be lost or rolled back accidentally. The outbox lets LakeCat retry projection from committed state. A drain acknowledges delivery only after every projection in the batch succeeds. Replay order is part of that contract: LakeCat selects undelivered outbox events by `created_at`, `event_id` in both embedded memory tests and the durable Turso store, and duplicate or malformed pending records fail with hash-only evidence before graph emission, lineage emission, or acknowledgement.

9.9 An Agentic QGLake Flow

The agentic path is the reason LakeCat has to be more than a passive catalog. Imagine a resilience supervisor agent investigating incidents:

1. The supervisor delegates table triage to a specialist agent.
2. The specialist asks LakeCat to plan a scan over `local.default.events`.
3. LakeCat resolves the agent identity and TypeDID context.
4. TypeSec authorizes the table scan and returns a restricted capability.
5. LakeCat narrows the projection and appends the required row predicate.
6. Sail plans against the current Iceberg metadata and delete manifests.
7. LakeCat returns governed plan and fetch-task responses.
8. The specialist summarizes only the allowed result shape.
9. LakeCat records scan and credential decisions into audit/outbox.
10. QueryGraph imports graph, policy, lineage, and bootstrap evidence.

The key point is the absence of raw storage reach. The specialist agent does not need broad cloud credentials to do its job. It needs a governed plan, a bounded task set, and a receipt trail.

The local fixture compresses this story into a short artifact-producing sequence:

```
cargo run -p lakecat-cli --features qglake-fixture -- qglake-fixture \
  --output target/qglake/lakecat-bootstrap.json \
  --drain-output target/qglake/lineage-drain.json \
  --principal did:example:agent
cargo run -p lakecat-cli -- qglake-verify-replay \
  --bundle target/qglake/lakecat-bootstrap.json \
  --drain target/qglake/lineage-drain.json \
  --principal did:example:agent
scripts/qglake-handoff-local.sh
cargo run -p lakecat-cli -- qglake-verify-handoff \
  --summary target/qglake-handoff/handoff-summary.json \
  --json
```

The fixture creates the sample table shape, installs a restricted policy, verifies governed scan planning, verifies fetch-scan-task reapplication, exercises delete manifest handling, probes credential-vend behavior for agents and trusted humans, verifies compact table commit-history evidence, exports QueryGraph bootstrap artifacts, drains the outbox, and proves the resulting bundle through QueryGraph's Rust verifier/importer. It is small, but it is not decorative. It is the acceptance story for a catalog that participates in the user workflow from notebook to agent.

10 Operating The Book's Example System

The local development posture is intentionally small:

```
cargo run -p lakecat-cli -- config
cargo run -p lakecat-cli -- storage-profile-list
cargo run -p lakecat-cli -- policy-list
cargo run -p lakecat-cli --features qglake-fixture -- qglake-fixture \
  --output target/qglake/lakecat-bootstrap.json \
  --drain-output target/qglake/lineage-drain.json \
  --principal did:example:agent
cargo run -p lakecat-cli -- qglake-verify-replay \
  --bundle target/qglake/lakecat-bootstrap.json \
  --drain target/qglake/lineage-drain.json \
  --principal did:example:agent
scripts/qglake-handoff-local.sh
cargo run -p lakecat-cli -- qglake-verify-handoff \
  --summary target/qglake-handoff/handoff-summary.json \
  --json
cargo run -p lakecat-cli -- bootstrap-export \
  --output lakecat-bootstrap.json
```

The important thing is what these commands exercise. They are not a separate product surface. They touch the same catalog, policy, bootstrap, and QueryGraph export contracts that the service uses.

11 Standards And Engine Boundary Decision Record

This chapter is the release decision record for the current catalog concepts. It answers three questions that come up whenever LakeCat vocabulary gets close to Iceberg vocabulary:

1. Which terms are standard Iceberg parlance?
2. Which terms are LakeCat, QueryGraph, TypeSec, Grust, or Sail extensions?
3. Which ideas are plausible future Iceberg-adjacent optional profiles?

The distinction matters because LakeCat should be ambitious without making ordinary Iceberg clients pay for that ambition. A PySpark job should be able to load and commit a table through normal REST catalog behavior. A governed agent should be able to receive a TypeSec-backed, Sail-planned, replayable proof of bounded work. QueryGraph should be able to import QGLake evidence, OpenLineage anchors, view receipt chains, credential posture, management proof, and commit proof. Those are three different audiences using one catalog foundation.

The implementation can be pictured as four concentric contracts:

standard Iceberg client contract
namespace, table, metadata location, optimistic commit

LakeCat catalog-control contract
identity, tenancy, Turso rows, CAS, idempotency, audit, outbox, replay

TypeSec and Sail governed-work contract
capability decision, restriction, receipt, projection, predicate, scan plan

QueryGraph semantic contract
QGLake handoff, OpenLineage, Croissant, CDIF, OSI, ODRL, graph, agents

The outer layers are additive. They can observe, prove, and compose catalog state, but they must not redefine the inner standard table contract.

11.1 Catalog Concepts In Plain Terms

This section names the catalog concepts directly, because the words can sound similar even when they live at different layers. The simplest way to read the architecture is this:

- Iceberg defines the portable table contract.
- LakeCat defines the durable catalog-control contract around that table.
- TypeSec defines the policy and capability evidence used by governed work.
- Sail defines the engine-grade interpretation of Iceberg table state.
- QueryGraph defines the semantic workflow built from accepted catalog proof.

The Rust service/catalog spine is a LakeCat implementation choice, not an Iceberg concept. Iceberg does not care whether a catalog is written in Rust, Java, Go, or Python. Iceberg cares that a client can ask for catalog config, list namespaces, load a table, and commit a new metadata pointer under ordinary REST catalog rules. LakeCat chooses Rust because the proof path should be short, typed, and close to the engine. In one process, LakeCat can parse the request, resolve identity, check tenancy, call TypeSec-style authorization, ask Sail for table interpretation, update

the `CatalogStore`, append audit and outbox rows, and return the standard response. The portable idea is not “Rust catalogs are standard.” The portable idea is that a catalog state transition can be deterministic, auditable, replayable, and bound to engine facts without extra indirection.

The Turso-backed local store is also a LakeCat implementation choice, not an Iceberg term. Iceberg needs durable catalog state and atomic metadata-pointer movement; it does not name Turso, SQLite, PostgreSQL, FoundationDB, or any other storage engine. LakeCat uses the Rust turso crate for the local durable spine because it gives the project a fast embeddable store while keeping the store contract explicit. The important abstraction is `CatalogStore`: tables, views, namespaces, policies, storage profiles, idempotency records, pointer logs, audit rows, outbox rows, and replay markers must be read and written in ways that bind decoded data back to the requested tenant, warehouse, namespace, and table. Turso is the current local vehicle. The future portable concept is stronger than a database choice: exact catalog retry, atomic pointer movement, row-binding checks, transactional event identity, and replayable durable proof.

The Iceberg REST namespace and table paths are the compatibility floor. These are standard Iceberg parlance: catalog config, namespace, table identifier, metadata location, table metadata, requirements, optimistic commit, snapshot, manifest list, manifest, data file, delete file, partition spec, sort order, and schema. A PySpark client should be able to use LakeCat as an Iceberg REST catalog without knowing what QGLake, TypeSec, Grust, or QueryGraph are. That is the line LakeCat must not cross. If normal table load or commit requires a non-standard QueryGraph endpoint, then LakeCat has stopped being an Iceberg-compatible catalog and has become a private application protocol. The LakeCat proof work should happen behind the route or beside the route, not in place of the route.

Commit CAS straddles standard language and LakeCat hardening. In Iceberg, a commit is optimistic: the client presents requirements that describe the table state it believes it is changing, and the catalog accepts the new metadata location only if the requirements still hold. LakeCat implements that as metadata-pointer compare-and-swap, which is standard in spirit: move from the previous metadata location to the new one only if the current state matches the expected state. LakeCat then adds proof around it: the idempotency record, metadata-object create-only posture, pointer log, audit row, outbox event, redacted conflict evidence, and replay admission verdict. CAS itself belongs to Iceberg-compatible commit correctness. The proof envelope is LakeCat’s catalog reliability layer, and a stripped-down form of it could become a future optional catalog reliability profile.

Idempotency, pointer logs, audit/outbox, and replay validation are not the normal vocabulary a PySpark user brings to Iceberg. They are LakeCat’s control-plane vocabulary. Idempotency means a retried commit with the same key must return the same accepted response or fail if the request drifted. Pointer logs mean operators and QueryGraph can reconstruct which metadata location was current after each accepted transition. Audit rows mean the catalog can explain who did what, under which principal and tenant, without exposing credentials or raw secrets. Outbox rows mean graph and lineage side effects are durable catalog facts that can be drained later rather than fragile request-path side effects. Replay validation means the drained event is checked again against a closed schema before Grust, OpenLineage, QGLake, or QueryGraph can inherit it. Those

are extensions today. They are also among the best future proposal candidates because they can be described without adopting LakeCat’s product names: “catalog event admission,” “exact retry,” “pointer history,” and “transactional lineage outbox” are useful neutral ideas.

Governed scan and credential paths are LakeCat plus TypeSec plus Sail. Iceberg itself says an engine can use table metadata to plan a scan. It does not say that a catalog must return a TypeSec receipt, a purpose string, a policy hash, an allowed-column list, a mandatory predicate, a credential TTL cap, or a redacted proof of storage scope. LakeCat introduces those as additive governed surfaces because agents and untrusted principals should not default to broad object-store credentials. In the governed scan path, TypeSec-style policy logic decides the capability and restriction; LakeCat binds that receipt to the current pointer, request identity, tenant, and audit context; Sail turns the restriction into engine work. In the governed credential path, raw credential vending is intentionally narrower: it is an audited exception for a principal that is allowed to read directly, and its response must leave redacted prefix hashes, storage-profile hashes, TTL evidence, and secret-ref posture rather than leaking credential material. This is an extension today. The future proposal candidate is not “TypeSec in Iceberg”; it is a policy-engine-neutral governed access profile with proof-carrying scan and credential posture.

QueryGraph and QGLake handoff surfaces are product integration surfaces, not standard Iceberg semantics. QueryGraph needs bootstrap evidence, management inventory, table and view manifests, OpenLineage event hashes, graph import hashes, view receipt chains, credential posture, storage-profile posture, commit-history proof, replay verdicts, and artifact hashes. LakeCat should emit those anchors because QueryGraph needs a trustworthy substrate. But ordinary Iceberg clients should not need to produce or consume them. The standardization candidate is narrower than the QueryGraph feature set: event identity, lineage binding, replay-admissible payloads, view lifecycle proof, and commit proof could be phrased neutrally. Croissant, CDIF, OSI, ODRL, TypeDID workflow composition, QGLake artifact layout, Grust graph import shape, and agent workflow reasoning should remain QueryGraph, TypeSec, and Grust architecture unless and until a narrower neutral profile emerges.

Typed Iceberg v4 support is different again. If Iceberg v4 introduces or solidifies table-format semantics, those semantics are Iceberg work and engine work, not LakeCat product invention. LakeCat can preserve unknown JSON fields as a compatibility bridge, but passthrough is not interpretation. The durable direction is for Sail to own typed support: metadata structures, manifest and delete semantics, scan planning, metadata-as-data, commit validation, and view/table behavior. LakeCat can then advertise an honest state: standard REST compatibility now, JSON preservation where necessary, and typed Sail-backed proof as those engine APIs mature.

That is the extension answer. LakeCat has extensions in the product sense, because it extends what a catalog can prove. Most of those extensions should not be proposed as mandatory Iceberg behavior. The ones worth considering for Iceberg-adjacent discussion are the parts that can be stated without proper nouns: exact retry, pointer history, redacted conflict proof, transactional catalog events, replay-admissible lineage, governed scan proof, governed credential posture, and view lifecycle proof. The proper nouns are where the implementation is being proven. The neutral concepts are what might someday be shared.

11.2 Decision Matrix

Concept	Iceberg meaning	LakeCat or Query-Graph meaning	Decision
Rust service/catalog spine	Iceberg does not specify an implementation language.	LakeCat keeps route handling, identity, tenancy, store transactions, Sail calls, TypeSec receipts, CAS, idempotency, audit, outbox, and replay close together in Rust.	Project implementation. The portable behavior is deterministic and replay-safe catalog state transition proof.
Turso-backed local store	Iceberg needs durable catalog state and atomic metadata-pointer movement, not a named database.	LakeCat uses the Rust turso crate as the durable local CatalogStore spine for tables, views, policies, storage profiles, idempotency, pointer logs, audit, outbox, and replay markers.	Project implementation. Generalize CAS, exact retry, pointer history, row binding, and transactional event identity, not Turso.
Iceberg REST namespace and table paths	Standard catalog compatibility: config, namespace operations, table load/create/drop, metadata locations, requirements, and optimistic commit.	LakeCat serves these paths while attaching server-side audit, outbox, and replay evidence.	Standard floor. Do not make normal clients depend on QueryGraph, TypeSec, Grust, QGLake, or LakeCat proof schemas.
Commit CAS	Optimistic metadata-pointer movement under commit requirements.	LakeCat binds the accepted transition to idempotency, pointer logs, audit, outbox, redacted conflict proof, and replay admission.	CAS is standard. The proof envelope is a future optional reliability-profile candidate.
Idempotency, pointer logs, audit/outbox, replay validation	Not generally standard table semantics, except insofar as they protect commit correctness.	LakeCat makes retries exact, pointer history inspectable, side effects recoverable, and down-serial provenance provably causal by attaching a causal dependency graph to each update. This is not a standard LakeCat extension.	LakeCat extension today. Strong candidate for neutral catalog reliability and event-admission proof schemas. By attaching a causal dependency graph to each update, it is a standard extension.
QGLake/Grust/TypeSec/QueryGraph lineage/bootstrap/management/view/commit proof	Not generally standard table semantics, except insofar as they protect commit correctness. TypeSec receipts.	LakeCat makes retries exact, pointer history inspectable, side effects recoverable, and down-serial provenance provably causal by attaching a causal dependency graph to each update. This is not a standard LakeCat extension.	LakeCat extension today. Strong candidate for neutral catalog reliability and event-admission proof schemas. By attaching a causal dependency graph to each update, it is a standard extension.

The standards answer is therefore not a blanket yes or no. LakeCat implements many extensions in the practical sense, because it extends what a catalog can prove. Only the parts that survive without LakeCat-specific names should be considered future Iceberg-adjacent proposals.

11.3 The Proper-Noun Test

The proper-noun test is the simplest review tool:

- If the concept requires the words LakeCat, QueryGraph, QGLake, TypeSec, Grust, Sail, or Turso to be meaningful, keep it as project architecture.
- If the concept can be described as behavior any Iceberg-compatible catalog or engine could implement, consider it a future optional-profile candidate.
- If the concept changes what a standard client must send or understand for an ordinary table load or commit, treat it as suspect until compatibility is proven.

That test keeps the proposal set small. Rust and Turso are excellent LakeCat choices, but they are not standards work. TypeSec receipts and QueryGraph import plans are essential product interfaces, but they should not become Iceberg requirements. Exact retry, pointer-history proof, redacted conflict proof, transactional event identity, replay-admissible evidence, governed scan proof, credential posture proof, view receipt chains, and lineage receipt binding are more plausible because they can be expressed without adopting the LakeCat stack.

11.4 Why The Engine Owns The Truth

The strongest technical reason to push work into Sail is that Iceberg table truth is not a set of easy strings. A governed proof that says “only these columns” is not reliable until it is mapped through Iceberg field ids, nested schema evolution, aliases, and projection rules. A proof that says “only these rows” is not reliable until the predicate has been interpreted by the same expression system that will plan the scan. A proof that says “these files” is not reliable until manifest metrics, residual predicates, partition transforms, delete files, sequence numbers, and snapshot context have been interpreted by the engine.

LakeCat can prove authority:

```
principal -> tenant -> warehouse -> namespace -> table
request -> TypeSec receipt -> current pointer -> CAS/idempotency
accepted state -> audit row -> outbox row -> replay admission
```

Sail should prove table work:

```
metadata -> schema field ids -> projection
metadata -> snapshots/manifests/statistics -> pruning
metadata -> equality/position deletes -> row visibility posture
metadata -> scan tasks / metadata-as-data / commit validation
```

When those two proofs are bound together, QueryGraph gets evidence that is both authorized and data-close. When LakeCat tries to do both jobs itself, it becomes a shadow engine: slower, less correct, and less reusable. Manifest metric decoding, delete planning, metadata-as-data, scan tasks, and typed v4 support should benefit Sail users, LakeCat governed reads, QueryGraph import, and the wider Rust lakehouse stack at the same time.

Sail is a particularly good engine boundary because it keeps the proof-heavy path Rust-to-Rust. LakeCat can receive a REST request, bind identity and tenancy, call TypeSec-style governance, ask Sail for typed table interpretation or a plan, commit through the Turso-backed CatalogStore, and persist replay proof without crossing a JVM sidecar, Python shim, or remote plugin boundary in the hot path. The public compatibility boundary remains Iceberg REST. The internal correctness boundary stays close to the Rust engine.

11.5 Workflow Consequences

For a PySpark user, the catalog remains ordinary. The user configures an Iceberg REST catalog, lists namespaces, loads a table, reads metadata, and commits with optimistic requirements. LakeCat's proof work stays server-side. The PySpark job does not need to understand TypeSec receipts, QGLake artifacts, Grust graph imports, or QueryGraph bootstrap.

For an operator, the catalog becomes inspectable. The operator can ask which principal moved a pointer, whether an idempotent retry drifted, which durable row was accepted, which outbox event is pending, which replay validator admitted the evidence, and whether downstream graph or OpenLineage projection came from accepted catalog state.

For a governed service, the catalog becomes enforceable. TypeSec narrows the request by purpose, capability, policy hash, allowed columns, mandatory predicate, TTL cap, and credential posture. LakeCat binds that decision to the current pointer and request identity. Sail turns the restriction into field-id projection, predicate binding, manifest/file pruning, delete posture, scan tasks, and plan evidence.

For an agent, the safest default becomes bounded work. Raw credentials are still possible for deliberately authorized principals, but they are audited exceptions. A restricted agent should usually receive Sail-planned tasks or a proof-backed fetch path whose scope, TTL, projection, predicate, and credential posture are already bound.

For QueryGraph, the catalog becomes a foundation rather than an application. LakeCat emits stable anchors: bootstrap hashes, management inventory hashes, policy-binding hashes, storage-profile hashes, commit-history hashes, view receipt-chain hashes, credential posture hashes, OpenLineage event hashes, graph/import hashes, artifact hashes, and replay verdicts. QueryGraph composes those anchors with Croissant, CDIF, OSI, ODRL, TypeDID, Grust graph state, and agent workflow state. LakeCat does not import QueryGraph. QueryGraph consumes LakeCat proof.

That separation is the release discipline. Standard clients get standard Iceberg behavior. Operators get durable Rust and Turso proof. Governed callers get TypeSec decisions bound to Sail engine work. QueryGraph gets a semantic handoff that is broad enough to build on and narrow enough not to fork Iceberg.

12 First Release Readiness

The first LakeCat release should not try to finish every idea in this book. It should release the catalog substrate that can already be proven locally. The right question is not “does LakeCat

contain the whole future QueryGraph stack?” It is “does LakeCat provide a compatible, durable, governed catalog foundation that QueryGraph can trust?”

For the first release, the release-blocking behavior is the catalog spine: standard Iceberg REST config, namespace, table-load, table-create, and table-commit paths; warehouse-prefixed routing; Rust service identity handling; the CatalogStore seam; the Turso-backed local store; memory-store parity for embedded tests; metadata-pointer CAS; idempotent commit replay; pointer logs; audit rows; outbox rows; and replay admission that rejects malformed durable evidence before graph or OpenLineage projection.

The governed path is also release-blocking because it is the reason LakeCat exists for agents. A restricted agent should be able to ask for a governed read, receive a TypeSec-style receipt, get Sail-planned work instead of broad storage authority, and leave behind replayable scan, fetch, credential, management, view, and commit-history evidence. Trusted raw credential vending can exist only as an audited exception with redacted storage-scope proof. Fetched scan replay treats the returned plan-task as evidence rather than arbitrary text: if it is present, it must be a non-empty LakeCat-issued token and it must not contain decorated location, query/fragment, or credential material before graph, OpenLineage, QGLake, or QueryGraph import can inherit the fetch proof. That keeps governed plan/fetch receipts about Sail-planned work, not a carrier for raw path or token claims. The same fetched replay path treats stats-fields as checked evidence when it is present: the array must be non-empty, duplicate-free, and bound to the effective stats fields before downstream proof can use it. The compact QGLake handoff proof now preserves that fetched side as its own fetchedRequestedStatsFields and fetchedEffectiveStatsFields evidence, so a handoff cannot prove only the planned stats narrowing while silently dropping what the fetch path actually returned. The local handoff script applies the same nonblank, duplicate-free array rule to projection, stats-field, and read-restriction allowed-column evidence before it writes that compact proof into the summary.

The QueryGraph handoff is release-blocking as an acceptance proof, not as a requirement for ordinary Iceberg clients. The local QGLake workflow must keep creating a bootstrap bundle, draining LakeCat lineage/outbox evidence, verifying replay, running QueryGraph verification/import, and saving a handoff summary whose artifact hashes, table/view counts, standards, OpenLineage receipts, graph hashes, policy anchors, credential proof, view receipt chains, and commit history agree. If that handoff cannot be reproduced locally, LakeCat may still serve an Iceberg endpoint, but it has not proven the QueryGraph foundation. The saved handoff verifier output must also repeat the lineage-drain catalogConfigProof; omission is treated as proof failure, not as an unspecified optional field. Extra fields inside that repeated proof are also rejected, so an archived verifier sidecar cannot append a new endpoint or compatibility claim beside the validated catalog configuration evidence. Archived handoff file paths are part of that proof. LakeCat resolves artifact paths under the handoff summary directory before hashing or parsing them, and the verifier rejects relative traversal outside that bundle for both hash verification and captured-output semantic reads. A matching hash from another directory is not accepted as QGLake evidence. The artifact manifest itself is also closed: the primary artifacts object, nested capturedOutputs object, and individual bundle, lineage-drain, QueryGraph import-plan, and

captured output artifact objects may carry only the path and SHA-256 evidence LakeCat checks. A saved summary cannot add an alternate hash, mirror artifact, or unverified capture beside an otherwise valid handoff bundle. The captured files themselves are closed at the root as well. Saved LakeCat replay output and QueryGraph verify/import output may carry only the fields LakeCat compares to the compact summary; a matching hash does not make an extra replay, QueryGraph, import, or application claim part of the proof. The same release gate treats raw view-lineage proof hashes as real digests, not placeholders: view replay receipts, tombstone view receipts, namespace receipt-chain hashes, and receipt-chain replay/OpenLineage hashes must be full SHA-256-shaped values before QGLake can archive them as accepted view proof.

The first release should explicitly defer work that belongs elsewhere or is not yet ready to claim. Typed Iceberg v4 semantics belong in Sail; LakeCat should advertise only the current JSON passthrough bridge with `typed-sail=unavailable` until Sail exposes stable typed support. Reusable graph taxonomy, traversal, Cypher, graph stores, and algorithms belong in Grust. Capability composition, TypeDID envelopes, secure-agent proof, and richer policy semantics belong in TypeSec. Coissant, CDIF, OSI, ODRL application composition, and agent-facing reasoning belong in QueryGraph. Cloud SDK secret managers beyond the current Vault and file-backed provider roots are future credential backends, not blockers for the catalog substrate.

As a working estimate, the first-release LakeCat catalog substrate is roughly 90 percent complete. That number is not a promise about the whole future QueryGraph architecture. It means the release-blocking LakeCat pieces are mostly present and locally proven: standard Iceberg REST namespace and table paths, the Rust service spine, the Turso-backed local store direction, metadata-pointer CAS, idempotency, pointer logs, audit and outbox rows, replay admission, governed scan and fetch proof, credential-vending receipt proof, management proof, view receipt chains, QueryGraph bootstrap, OpenLineage replay, and QGLake handoff/import evidence.

The remaining 5-10 percent is concentrated in release engineering and dependency-boundary cleanup rather than a new conceptual layer. The broad local gate has already been recorded from the current handoff path, including QueryGraph verification and import under `--locked`; the release task is to keep that gate green after each dependency-boundary change and rerun it from the final release commit. The Sail dependency is release-explicit: LakeCat builds Sail from a Cargo git dependency on the `lakecat` branch of `github.com/querygraph/sail` until the required Sail APIs are published. The Grust contract is likewise explicit: LakeCat and QueryGraph follow the active local Grust 0.10 path checkout, and LakeCat binds Turso-backed graph projection to the dedicated `grust-turso` crate. `grust-local` keeps fast memory-backed projection tests, while `grust-turso-local` proves durable graph projection through a bootstrapped `grust_turso::TursoGraphStore`. QueryGraph's `qg-rust` checkout uses the same local Grust path for `lakecat-verify` and `lakecat-import`, keeping the end-to-end graph import path aligned with the catalog graph sink. LakeCat's own graph responsibility stops at the catalog-facing envelope: graph sinks reject blank projection identity, non-object properties, and table-scoped labels without table identity before handing anything to Grust. Persistence, traversal, Cypher, and graph mutation semantics stay in Grust. The service also redacts Grust Turso graph-sink connect/bootstrap failures to graph-

store-path-hash and backend-error-hash evidence so release logs do not capture raw graph database paths.

README, status, changelog, book artifacts, and version notes must be refreshed from the same clean proof commit. The already-published v0.1.0 tag should not move; current post-tag hardening stays under Unreleased while the workspace version remains 0.1.0. For a future version-bump release, tag only after the broad local gate, QGLake handoff, QueryGraph locked verify/import, dependency-contract check, and book validation all pass together. Tracked book artifacts are refreshed deliberately with docs/book/build.sh; the release-candidate gate validates a fresh EPUB/PDF/MOBI build in a temporary LAKECAT_BOOK_DIST_DIR and gives Calibre a temporary CALIBRE_CONFIG_DIRECTORY, so neither binary artifact metadata nor converter preference state can dirty the candidate commit or operator profile.

Manual GitHub Actions is deliberately narrower than the local release proof. When an operator triggers it, the workflow runs the local dependency contract, the workflow-trigger self-test, and the release-version contract before the Rust matrix. It does not own release-proof freshness. That check belongs to the clean local release-candidate gate, because ordinary post-proof code hardening should report “the old proof is stale” locally rather than create a surprising failing cloud run. In practice: use manual CI as a second environment after the local gates are green, not as the source of release truth.

That leaves important work after the first release, but it should stay out of the first-release blocker list unless the scope changes. Typed Iceberg v4 semantics belong in Sail. Cloud SDK-backed secret resolvers belong behind the existing TypeSec-gated provider seam. Reusable graph taxonomy, traversal, stores, and Cypher behavior belong in Grust. Full Croissant, CDIF, OSI, ODRL application composition and agentic workflow semantics belong in QueryGraph and TypeSec above LakeCat. The first release should prove the catalog foundation, not absorb every future system.

The release evidence is concrete:

```
scripts/check-release-readiness.sh --release-candidate
scripts/qglake-handoff-local.sh
docs/book/build.sh
scripts/check-book-artifact-contract.sh docs/book/dist
scripts/check-local-dependency-contract.sh
```

The current full local release-candidate proof was refreshed on June 26, 2026 from clean head 6f5c1739. It passed with tracked book artifact validation, the checked-in release-proof contract in clean candidate mode, the strengthened post-tag release-posture contract for the published v0.1.0 baseline, the querygraph/sail lakecat git-dependency source assertions, temporary book build, executable book artifact contract, QueryGraph locked verify/import, Grust Turso graph projection proof, bundle hash sha256:88e38f620068d13cb14cb3ad3f102558b694482a87b45f09c08419ed93cf17cb, graph hash sha256:7c6aa85c544d67953edf7bd168a85d8cfaa87a2f48f2732b77cf443031db01a7, OpenLineage hash sha256:c86ce5e6a82ad241a67a99301e802b42c9f07b020869b3893103e9b780561aab, QueryGraph import hash

sha256:8c662182623d7c51bc1397ffffd8228c1c73c82130c4bb42f5ca9d1e08b4e220, and the final clean-tree check.

LakeCat also carries a smaller proof-freshness contract for the release docs themselves:

```
scripts/check-release-proof-contract.sh
```

That command verifies the active docs agree on the latest full release-candidate proof commit, that the proof commit is an ancestor of the current tree, and that any commits after it are limited to documentation and checked-in book artifact refresh. This avoids an infinite proof-update loop: the heavy gate proves the candidate, then a narrow documentation commit records the proof. If Rust code, manifests, workflows, dependency bridges, release scripts, or other executable behavior changes after the cited proof commit, the old proof is no longer enough and the full release-candidate gate must run again. The quick and ordinary full release-readiness gates surface that distinction without becoming release proof themselves. They print a non-failing freshness note when executable changes exist after the recorded proof commit, so a narrow development slice can still finish while the operator sees that final release proof requires a new `scripts/check-release-readiness.sh --release-candidate` run.

The contract requires a clean tree by default. While editing the proof contract or release docs, maintainers can run:

```
LAKECAT_RELEASE_PROOF_ALLOW_DIRTY=1 scripts/check-release-proof-contract.sh
```

That is only a local self-test mode. It still checks unstaged, staged, and untracked paths against the same post-proof allowlist, so a dirty Rust source file or manifest does not masquerade as documentation-only release evidence.

The full release-candidate gate runs the same contract in candidate mode:

```
LAKECAT_RELEASE_PROOF_CANDIDATE=1 scripts/check-release-proof-contract.sh
```

That mode is intentionally narrow. It still requires a clean tree and coherent active proof references, but it allows the current clean HEAD to become the next proof commit. After the heavy gate passes, the follow-up documentation and book artifact refresh records that new proof commit.

The already-published `v0.1.0` tag is a baseline, not something to move. The already-published `v0.2.0` tag is a baseline, not something to move either. While the workspace version remains `0.1.0` and HEAD is past `v0.1.0`, post-tag hardening stays under `Unreleased`. The release version contract checks that shape directly so a follow-up proof commit cannot accidentally look like a second same-version release. It also requires the release checklist to scope tagging chores to a future version-bump release, not the already-published baseline. Finally, it derives the expected versioned changelog heading date from the existing tag, not the current day, so the published `0.1.0 - 2026-06-23` release heading remains stable while hardening continues under `Unreleased`.

The quick check is acceptable while landing a narrow slice:

```
scripts/check-release-readiness.sh --quick
```

The quick check is not a release claim, but it does validate the tracked docs/book/dist artifact contract so stale EPUB/PDF/MOBI deliverables are caught before a full release-candidate build regenerates temporary artifacts. The full gate is the first-release claim. It runs local dependency-contract checks, workflow-trigger checks, formatting, default workspace tests, feature matrix tests, Turso rows, Sail/TypeSec/Grust integration rows, all-features workspace tests, an out-of-tree book build with artifact validators, and the QGLake handoff proof. Cloud CI remains manual-only until that local proof is boring. A release should be cut from the state those commands prove, not from a hope that a remote runner will discover the truth later.

The handoff portion must stay locked in QueryGraph as well as in LakeCat. The local gate intentionally runs QueryGraph verification and import with Cargo's `--locked` mode so a stale QueryGraph dependency graph is a local failure, not a cloud surprise. When Grust or TypeSec publish new crates, the right sequence is to refresh the QueryGraph lockfile locally, prove `lakecat-verify`, `lakecat-import`, and QueryGraph tests under `--locked`, then rerun LakeCat's full release gate. That makes the QueryGraph import contract part of the release artifact instead of a best-effort downstream hope.

The current gate also proves a subtle QGLake contract that is easy to miss. The handoff readiness probe reads `GET /catalog/v1/config`, and that read is a catalog event. It must therefore carry the same QGLake agent identity as the rest of the handoff, otherwise a harmless-looking readiness curl becomes an anonymous `catalog.config-read` proof in the drained evidence. The handoff script binds that config-read proof to the agent, writes a canonical handoff-verifier output artifact, hashes it into the summary before strict verification, and keeps management list proof, policy-upsert proof, and storage-profile upsert proof in the distinct sections the verifier compares. That is the local substitute for cloud hope: the artifact graph is checked before the summary is accepted.

For this release, replay admission is part of the local proof rather than a best-effort projection filter. The service closes event payloads before graph, OpenLineage, QGLake, or QueryGraph import can inherit them; the compact handoff verifier then repeats the same checks against saved artifacts. Appendix A lists the replay-admission contract, Appendix B captures credential redaction, Appendix C captures view receipt chains, and Appendix D maps the QGLake handoff proof.

13 Appendix A: Replay Admission Matrix

Replay admission is the line between durable catalog fact and downstream proof. An outbox row may be stored, but it is not allowed to become graph, OpenLineage, QGLake, or QueryGraph evidence until LakeCat has checked its shape, scope, authority, and redaction posture. The matrix below keeps the best review version of those checks in one place.

Event family	Must prove	Rejects	Why it matters
catalog.config-read	Warehouse scope, v4 bridge defaults, overrides, advertised standard REST endpoints, governed plan/fetch/credential endpoints, bootstrap and lineage-drain endpoints, authorization receipt	Duplicate config keys, unsupported lakecat.format.v4* claims, missing standard or governed endpoints, extra unverified config fields	QueryGraph imports the same compatibility contract LakeCat advertised, without turning config into a free-form claims bag
Namespace lifecycle	Warehouse, namespace path or components, matching lifecycle action, valid principal	Malformed namespace paths, count drift, action drift, unclosed wrapper fields	Standard catalog discovery stays attributable and replayable
Table lifecycle	Root table identity, warehouse/namespace/table hints matching that identity, action receipt, soft-delete format evidence when present	Scope drift, action drift, conflicting format-version/format_version aliases, undecorated-location violations	Table proof follows the Iceberg object it claims to describe
Table commit	Table scope, authorization receipt, pointer transition, principal, idempotency/request/response/policy hashes, format, snapshot, timestamp	Credential-bearing or decorated metadata-location evidence, unclosed nested commit fields, forged event type	Commit proof stays tied to the accepted pointer movement without leaking object paths
Governed scan and fetch	Principal, table-plan-scan action, requested/effective projections, requested/effective stats fields, required filters, matching top-level and receipt read restrictions, non-empty full policy hashes, past is more precise than future	Actorless scan proof, action drift, missing or widening projection/filter evidence, placeholder policy hashes, non-LakeCat or credential-bearing plan-task values	A saved scan proof shows the same TypeSec restriction Sail planned and LakeCat replayed
QueryGraph lifecycle	Warehouse scope, v4 bridge defaults, overrides, advertised standard REST endpoints, governed plan/fetch/credential endpoints, bootstrap and lineage-drain endpoints, authorization receipt	Duplicate config keys, unsupported lakecat.format.v4* claims, missing standard or governed endpoints, extra unverified config fields	QueryGraph imports the same compatibility contract LakeCat advertised, without turning config into a free-form claims bag

The pattern is deliberately repetitive in code but not in the narrative: every event must be scoped, actor-bound, action-bound, closed over the fields LakeCat actually verifies, and redacted before it becomes downstream evidence.

14 Appendix B: Credential And Redaction Contract

Credential handling is where catalog convenience can become security risk. LakeCat's posture is therefore narrow: governed Sail-planned reads are the default for agents, while raw credentials are explicit audited exceptions.

Surface	Accepted shape	Redaction rule	Failure posture
Warehouse storage root	Plain URI rooted in the warehouse; no query, fragment, userinfo, or dot traversal	Return warehouse-storage-root-hash rather than raw root on errors	Reject before memory or Turso persistence
Storage-profile prefix	Provider-compatible URI prefix, longest-prefix match, no ambiguous tie	Return storage-profile-prefix-hash or location-prefix-hash	Fail closed rather than guessing a credential root
Public config	Non-secret string hints such as region, endpoint labels, or purpose	Return public-config-key-hash for rejected keys/values	Reject secret-looking keys, values, and LakeCat-reserved proof keys
Secret reference	Clean external locator such as vault://, aws-sm://, gcp-sm://, or azure-kv://	Store and replay secret-ref-present, provider, and full secret-ref-hash	Reject decorated URIs, unsupported schemes, traversal, userinfo, query, fragment
Resolver dispatch	TypeSec authorizes the exact secret-ref resource before lookup	Resolver failures report provider plus secret/error hashes	Denied decisions do not call the resolver
Returned credentials	Prefix remains inside the selected profile; TTL honors policy cap; LakeCat-owned evidence is canonical	Replay prefix hashes and issuer-config hashes, never raw secret material	Reject widened prefixes before recording credentials.vend-attempted
Restricted agent branch	Zero credentials, explicit block reason, raw exception false, TTL cap present	Empty credential-response evidence is explicit proof, not absence	Agent is steered to governed Sail-planned reads
Trusted human branch	Positive credential count only when policy allows raw exception	Returned prefixes are full hashes; exception reason is recorded	Raw credentials remain auditable exceptions, not the default path

The important distinction is that storage-profile management proves a credential root exists, while credential vending proves whether a caller received raw storage reach. QueryGraph can

reason over both without seeing the raw root, secret URI, backend error, token, or returned credential material.

15 Appendix C: View Receipt-Chain Rules

LakeCat's first release does not claim full Iceberg view-history semantics. It keeps a compact, catalog-owned receipt chain so QueryGraph can prove which view version was exported, updated, dropped, or tombstoned. Richer view semantics remain a Sail-aligned future target.

Rule	Required evidence	Rejects
Stable identity	lakecat:view:<warehouse>:<namespace>Stable IDs that do not match plus matching component warehouse/namespace/name fields	Stable IDs that do not match warehouse/namespace/name fields
First receipt	Version 1 upsert, no previous version, no previous receipt hash	Zero-version chains, first tombstones, forged first previous-link fields
Later upsert	Previous receipt hash points at the prior receipt; version advances by exactly one	Skipped versions, unsupported operations, broken previous links
Drop tombstone	Operation drop, previous receipt hash set, view-version remains the deleted version	Drops that advance the version or omit the previous receipt
Recreate	New upsert links after the tombstone and advances to the next version	A second unrelated version-1 chain for the same stable view id
Chain header	receipt-count, latest version, latest operation, and tombstone flag agree with the last receipt	Headers that shed receipts or present a forged head
Hash coverage	receipt-hashes, drop-receipt-hashes, and chain-hashes exactly cover the nested receipt bodies and chains	Convenient full-looking hashes that do not match the embedded chain
Durable row binding	Turso row warehouse/name-space/view columns agree with decoded receipt JSON	A valid-looking receipt body riding a corrupted row index into another view
QueryGraph handoff	Accepted view receipt hashes match bootstrap view artifacts and compact handoff proof	Spliced receipt arrays from another view or another run

The receipt chain is intentionally compact: it proves catalog state changes and tombstones, not full SQL engine view semantics. That is enough for QueryGraph to reject a bootstrap that lost or borrowed view proof while LakeCat keeps richer view interpretation out of the catalog.

16 Appendix D: QGLake Handoff Proof Map

The QGLake handoff is a compact acceptance artifact over several raw artifacts: the LakeCat bootstrap bundle, lineage-drain response, QueryGraph import plan, captured verifier output, captured replay output, and service log. The summary is useful only because every compact field can be traced back to a parsed raw artifact and a replay-admitted catalog event.

Proof section	Binds	Compared against
requestIdentityProof	Principal subject/kind, identity source/state, lineage-drain receipt, TypeDID hash slots	Captured LakeCat replay request identity and saved handoff-verifier output
queryGraphBootstrapProof	Bundle hash, graph hash, OpenLineage hash, QueryGraph import hash, standards, artifact counts, policy count, view receipt hashes	Bootstrap replay event, bundle manifest, QueryGraph verify/import output
governedScanProof	Planned/fetched task counts, projections, stats fields, filters, read restrictions, policy hashes, TTL cap	Scan replay sections and operator replay text
tableCommitHistoryProof	Commit count, positive increasing sequence numbers, duplicate-free commit hashes, replay/OpenLineage hashes	Pointer-log read replay and compact summary
viewReceiptChainProof	Accepted view versions, receipt hashes, receipt-chain hashes, tombstone receipts, verified-chain counts	View replay sections, bootstrap view artifacts, receipt-chain reads
managementProof	Server/project/warehouse/policy/storage-profile counts, ID arrays, graph event counts, replay/OpenLineage hashes	Management-list replay sections and QueryGraph bootstrap policy count
policyUpsertProof	Policy id, ODRL hash, principal/action, graph events, replay/OpenLineage hashes	Policy upsert replay and policy-list proof
storageProfileUpsertProof	Profile id, provider, issuance mode, location-prefix hash, secret-ref posture, principal/action	Storage-profile upsert replay and credential branches
credentialVendingProof	Restricted-agent block, trusted-human exception, credential counts, prefix hashes, TTL caps, storage-profile anchor	Credential replay sections and operator replay text
Artifact manifest	Bundle, lineage drain, import plan, captured outputs, service log, saved handoff verifier output	Recomputed file hashes under the handoff summary directory

The verifier enforces three global rules. First, artifact paths resolve under the handoff summary directory before hashing or parsing, so a summary cannot splice matching bytes from outside the bundle. Second, all integrity anchors are full sha256:-prefixed 64-hex digests, not readable placeholders. Third, saved sidecars are closed over the schema LakeCat compares; a hash match does not make extra unverified fields part of the proof.

This is why the handoff is a release acceptance artifact instead of a convenient log bundle. It proves that LakeCat replay, graph projection, OpenLineage, QueryGraph verification, and QueryGraph import all agreed on the same catalog state.

17 Future Work

LakeCat is a direction more than a single release. The next stage should keep the architecture more true without pushing engine, graph, or security semantics back into the catalog.

1. Point the Sail git dependency at upstream (or a published crate) only after the required helpers land there, then retire the querygraph/sail lakecat branch and its dependency-contract assertions.
2. Keep pushing reusable table-status, planning, metadata-as-data, manifest, delete, and commit helpers into Sail.
3. Treat typed Iceberg v4 as a Sail project after formal specification adoption. The active v4 draft includes relative metadata locations, but LakeCat should not turn a moving draft into catalog behavior; its bridge remains explicit and JSON-only until Sail can model and plan the adopted semantics.
4. Make the transactional outbox the only path for graph and lineage side effects, while keeping graph persistence, traversal, and Cypher in Grust.
5. Expand TypeSec-backed capability checks until every privileged path is unbypassable, with raw credentials remaining an explicit audited exception.
6. Keep OSI, Croissant, CDIF, ODRL composition, and agent workflows in QueryGraph. Before the next public release, refresh QueryGraph's older dependency-guide examples to match its live Grust 0.10.0 and TypeSec 0.8.0 runtime manifest; its LakeCat importer already validates the current receipt-chain and Grust graph handoff contract.
7. Prove the bootstrap bundle through QueryGraph import on every meaningful public-surface change, and keep local release evidence ahead of cloud CI.

The payoff is a catalog foundation that feels ordinary to Iceberg clients and rich to QueryGraph. Tables remain portable. Policies become enforceable. Graph and lineage become replayable. Sail plans close to the data. QueryGraph gets a trustworthy substrate for the next version of its lakehouse intelligence.