

LAKECAT

OCELOT: GOVERNED ICEBERG REST WITH
PROOF BUILT IN



ALEXY KHRABROV



LakeCat

Ocelot: Governed Iceberg REST with Proof Built In

Alexy Khrabrov

2026-08-26

Contents

1	Preface	2
2	Catalogs, Iceberg, and What LakeCat Adds	3
2.1	What a catalog is	3
2.2	What Iceberg makes the catalog responsible for	4
2.3	What LakeCat adds without changing Iceberg	5
3	The Vocabulary	5
4	The Boundary Model	8
5	Where LakeCat Stands Today	10
5.1	The Ocelot baseline	10
6	Why the Catalog Stays Thin	12
7	The Architecture	13
7.1	The CatalogStore seam	14
7.2	The read path	14
7.3	The commit path	15
7.4	The durable spine	17
8	The Benchmark Suite	17
8.1	Measuring catalogs, not object stores	18
8.2	Where LakeCat lands	18
8.3	What it took to get there	18
8.4	Audit and idempotency: the durable cost of features	19
8.5	The object-store cache and the scan benchmark	19
8.6	Rust versus the JVM	20
8.7	The read-write benchmark and the stock-client round-trip	20
9	The Siblings and the Engine Path	21
9.1	Grust: the graph boundary	21
9.2	TypeSec: the authorization boundary	22
9.3	Sail and the v3→v4 path	22
9.4	The semantic handoff: Croissant, CDIF, OSI, OpenLineage	24
9.5	The full stack	25
9.6	Ocelot's dependency baseline	27
9.7	Implementation shape	27
10	Workflow Examples	28

10.1	Starting The Catalog	28
10.2	Registering The Warehouse Shape	29
10.3	Storage Profiles And Credential Roots	30
10.4	A PySpark User Reads Iceberg	31
10.5	A Notebook Requests Credentials	32
10.6	A View Becomes Part Of The Catalog Story	32
10.7	QueryGraph Bootstrap	34
10.8	Draining The Outbox	35
10.9	An Agentic QGLake Flow	35
11	Operating The Book’s Example System	37
12	Standards And Engine Boundary Decision Record	37
12.1	Catalog Concepts In Plain Terms	38
12.2	Decision Matrix	41
12.3	The Proper-Noun Test	43
12.4	Why The Engine Owns The Truth	43
12.5	Workflow Consequences	44
13	Ocelot: The 0.3 Release	44
13.1	What shipped	45
13.2	Stock-client conformance	45
13.3	Governed reads and bounded credentials	45
13.4	The handoff as release proof	46
13.5	Dependency provenance	46
13.6	What Ocelot does not claim	47
13.7	Release and announcement posture	47
13.8	Reproducing the proof	47
14	Appendix A: Replay Admission Matrix	49
15	Appendix B: Credential And Redaction Contract	51
16	Appendix C: View Receipt-Chain Rules	53
17	Appendix D: QGLake Handoff Proof Map	55
18	Future Work	57

1 Preface

LakeCat is a Rust-native, Iceberg-compatible catalog foundation for QueryGraph. It begins from one conservative claim and one ambitious one. The conservative claim: an ordinary Iceberg REST catalog must keep working for ordinary engines — Spark, Flink, Trino, DuckDB, PyIceberg, Sail — with no new protocol to learn. The ambitious claim: the *next* catalog also has to be a governed control plane for Rust-first planning, semantic graph handoff, lineage, and agent access.

LakeCat holds both at once by drawing a sharp line. It keeps Iceberg compatibility at the boundary and moves the work that needs deep table-format knowledge to the engine that can reason about it. In this repository the catalog is **LakeCat** and the engine-side lakehouse implementation is **Sail**. Graph behavior belongs to **Grust**; governance and capability proof belong to **TypeSec**; the end-to-end integration target is **QueryGraph**.

This edition is aligned with **LakeCat 0.3.0, Ocelot**, published on July 4, 2026. Ocelot is the point where the architecture in this book became a shipped, announcement-ready claim: a stock PyIceberg client completed a real REST round-trip, the release-candidate gate passed from a clean tree, and the QGLake handoff proved that LakeCat replay, OpenLineage evidence, Grust projection, and QueryGraph verification/import agreed on the same catalog transition.

This book builds from first principles. Chapter 2 explains what a catalog is, what Iceberg makes the catalog responsible for, and what LakeCat adds without changing the table format. Chapter 3 fixes the vocabulary — the single hardest thing in a layered system — so every later term has an owner. Chapter 4 draws the boundary model: which repository owns which concept, and which ideas might someday become neutral, standardizable profiles. The middle chapters walk the live architecture: the service spine, the read path, the commit path, and the durable store. The final chapters cover the engine boundary and the v3→v4 path, the graph/security/lineage handoffs, the QueryGraph/QGLake acceptance flow, worked examples, and the Ocelot release proof and its deliberately bounded scope.

Each concept is defined once and then used. Where a later chapter needs a term, it links back rather than restating it.

2 Catalogs, Iceberg, and What LakeCat Adds

2.1 What a catalog is

A data catalog is often described as a place that lists datasets. That is true but too small. A real catalog is the control plane between names, storage, metadata, identity, and intent. At minimum it answers four questions:

1. What table does this name mean?
2. Where is its current metadata?
3. Who may read, write, plan, or administer it?
4. What changed, when, and under whose authority?

In a traditional database the catalog is embedded in the engine: one system owns table definitions, statistics, permissions, and the transaction log. A lakehouse pulls that apart. Data files live in object storage, metadata files live beside them, and several engines read and write the same tables. The catalog becomes the agreement point — it maps a logical table name to the current metadata pointer and arbitrates updates to that pointer.




Figure 1: Diagram 1

That pointer is deceptively important. If the catalog points at metadata version 17, the table is version 17. If a writer prepares version 18 and wins the compare-and-swap, the table becomes 18; if it loses, nothing partially changes. The catalog is not the table format, but it is where table history becomes visible and durable. For human analytics this sounds like bookkeeping. For agentic systems it is a trust boundary: a catalog can know the principal, warehouse, namespace, table, snapshot, requested columns, row restriction, storage profile, and policy receipt — and if it captures that before planning and commits it after state changes, it becomes a governed control plane rather than a passive address book.

2.2 What Iceberg makes the catalog responsible for

Apache Iceberg is a table format for large analytic tables. Its core idea is to put the table’s truth in explicit metadata files so engines can plan reads and validate writes without directory listing or fragile storage conventions. A current metadata file names schemas, partition specs, sort orders, snapshots, and properties; snapshots point to manifest lists; manifest lists point to manifests; manifests describe data and delete files.

The catalog’s role in Iceberg is intentionally narrow. Standard clients must be able to load table metadata, create namespaces and tables, commit changes, and sometimes receive credentials or scan tasks — over a documented REST shape:

```

GET /v1/config
GET /v1/{prefix}/namespaces
POST /v1/{prefix}/namespaces/{ns}/tables
GET /v1/{prefix}/namespaces/{ns}/tables/{table}
POST /v1/{prefix}/namespaces/{ns}/tables/{table}    # commit: requirements +
updates

```

The rule that governs everything else in this book: **a normal client must never have to call a non-standard endpoint to read an ordinary table.** If it does, compatibility is already broken.

2.3 What LakeCat adds without changing Iceberg

LakeCat’s promise is *compatibility first, evidence second, semantics above the catalog*. A Spark or PyIceberg client sees an Iceberg REST catalog. A QueryGraph or governed-agent client may ask for richer proof. Both share the same table because the portable truth stays in Iceberg metadata and the extra evidence sits beside it — never inside the table format.




Figure 2: Diagram 2

Iceberg metadata stays pristine. Policy, graph, lineage, and agent state are *derived* control-plane or graph data. The table remains an Iceberg table.

3 The Vocabulary

A layered system fails the moment its words blur. The most common mistake is to call every useful LakeCat feature an “Iceberg extension,” which makes the standard boundary too large. The test is simple: **ask what breaks if a client knows nothing about LakeCat.** If a PySpark job cannot load, commit, or drop a table without the concept, it belongs to standard compatibility. If PySpark keeps working but operators, governed agents, or QueryGraph gain stronger evidence, the concept is an additive surface.

Every term in this book falls into one of six categories.




Figure 3: Diagram 3

Standard Iceberg parlance — the words Iceberg already owns: catalog, namespace, table identifier, current metadata location, snapshot, manifest list, manifest, data file, delete file, schema and partition evolution, optimistic commit, REST compatibility. LakeCat must implement these faithfully. Changing their meaning means losing compatibility.

LakeCat implementation — how this Rust catalog satisfies the contract reliably: the service spine, the CatalogStore trait, the Turso-backed durable store, normalized idempotency rows, pointer logs, audit rows, outbox rows, redaction rules, and replay validators. These make ordinary Iceberg behavior atomic, inspectable, and replayable. They are not extensions; they are good engineering behind a standard surface.

LakeCat optional extensions — additive APIs beside the standard path: management inventory, commit-history inspection, view proof, credential-root posture, replay verification, OpenLineage projection, and QueryGraph/QGLake bootstrap bundles. They help operators, agents, and QueryGraph without becoming hidden requirements for ordinary table access.

TypeSec governance — authority and receipt semantics: who may act, capability proof, TypeDID envelopes, ODRL-derived restrictions, and agent posture. TypeSec *decides*; LakeCat *carries and records* the decision.

QueryGraph semantics — composition above the catalog: Croissant, CDIF, OSI, ODRL, OpenLineage, and Grust graph projections built from the governed source of truth.

Sail engine truth — table-format interpretation, metadata-as-data, scan planning, commit validation, and typed v4 behavior. The catalog binds its proof to engine truth rather than reimplementing the format.

The single reference below replaces the per-claim restatements that earlier drafts repeated; read each LakeCat claim through its category.

Claim	Standard reading	Category	Portable idea (if any)
Rust service / catalog spine	Iceberg needs a catalog authority, not a language	LakeCat implementation	“A catalog can prove what it committed, planned, vended, and emitted”
Turso-backed durable store	Iceberg needs durable state + atomic pointer movement, not a named DB	LakeCat implementation	CAS, exact-retry, row/content binding, pointer-history proof
REST namespace/table routes + commit CAS	Standard compatibility	Standard Iceberg	— (must stay ordinary)
Idempotency, pointer logs, audit/outbox, replay validation	Mostly outside the table contract	LakeCat implementation / extension	Stable catalog-event identity; scoped replay admission
Governed scan with policy receipt	Iceberg gives scan inputs, not receipts	TypeSec governance	Engine-proved effective projection + predicate
Credential vending posture	Catalog-adjacent, not ordinary table semantics	TypeSec governance	Bounded, redacted, engine-neutral credential profile
QueryGraph / QGLake / OpenLineage handoff	Not required for table access	QueryGraph semantics	Redacted, replayable lineage/view/commit evidence
Typed Iceberg v4 behavior	Belongs to the format as it evolves	Sail engine truth	Engine-owned v4 interpretation, not catalog JSON parsing

4 The Boundary Model

The vocabulary tells you which category a word lives in. The boundary model says which *repository* owns the work and why. LakeCat is deliberately thin: it keeps identity, tenancy, metadata-pointer state, policy gates, idempotent commits, and integration events — and delegates everything reusable.




Figure 4: Diagram 4

The ownership rule, stated once:

Concern	Owner	LakeCat keeps only
Iceberg format, manifests, scan planning, pruning, delete handling, v4	Sail	The call into Sail and the proof binding its result
Graph schema, taxonomy, traversal, stores, Cypher	Grust	The catalog-facing sink/projection boundary
Authorization, policy composition, capabilities, TypeDID, credential decisions	TypeSec	The request for a decision and the persisted receipt
Croissant/CDIF/OSI/ODRL/OpenLineage composition, agent workflows	QueryGraph	The governed bootstrap bundle it emits
Identity, tenancy, pointer CAS, idempotency, audit, outbox, replay	LakeCat	All of it — this is the thin catalog

This is also the answer to the recurring “is it an extension or a standard?” question. LakeCat is opinionated in code and modest in standardization. “Use Rust,” “use Turso,” “use TypeSec,” “import QGLake,” “project into Grust” are project choices, not proposals. The portable ideas are narrower and stated without product names: *reject idempotency drift*; *record redacted pointer history*; *emit transactional catalog-event identity*; *admit only scoped replay evidence*; *prove a governed scan was narrowed by an engine*. A proposal that forced every Iceberg catalog to understand QueryGraph, TypeSec, Grust, or Turso would narrow the ecosystem; a proposal that says “a catalog may publish redacted, replayable commit evidence” leaves room for many implementations. Prove the stronger shape locally, then extract only the database-neutral, policy-neutral, engine-neutral part.

5 Where LakeCat Stands Today

LakeCat is not only a design. The implementation already has a Rust service/catalog spine, a Turso-backed durable store, Iceberg REST-compatible namespace and table paths, hardened commit and replay evidence, governed scan and credential proof, and QueryGraph/QGLake handoff surfaces. Read each of those through its category (Chapter 3) rather than flattening them into “Iceberg extensions”: some are ordinary catalog implementation, some are optional LakeCat APIs, some are TypeSec governance proof, and only a few are plausible future standardization candidates.

5.1 The Ocelot baseline

The published baseline is v0.3.0, **Ocelot**. It turns the phrase “compatible, durable, governed catalog” into four separately checkable claims:

Claim	Ocelot evidence
Stock-client compatibility	An unmodified PyIceberg RestCatalog can create a namespace and table, append real data, list the table, load it, and read it back. REST map fields are objects, error types follow Iceberg’s ErrorModel, and namespace behavior matches stock-client expectations.
Durable catalog state	Turso transactions bind pointer compare-and-swap, pointer history, audit, outbox, and idempotency. An exact retry replays; a conflicting retry is rejected.
Governed work	TypeSec restrictions are applied before Sail plans. Raw credential vending remains an explicit, redacted, audited exception.
Verifiable handoff	The QGLake gate creates a bootstrap bundle, drains lineage, projects through Grust, and runs QueryGraph’s locked verifier and importer over hash-bound artifacts.

The release does not put those extra proofs in front of ordinary clients. The standard REST surface remains the floor; governance, replay, lineage, graph, and QueryGraph bootstrap are additive paths beside it.

The most important word is **beside**. New capability sits beside the Iceberg REST path, never in front of it. A standard client never has to present a TypeDID, parse a QueryGraph bundle, read a Grust edge, or inspect an OpenLineage receipt to load a normal table. So the same catalog shows a different face to each caller:

Caller	What they see
PySpark / PyIceberg user	Ordinary Iceberg REST. Configure a catalog, create a namespace, write and load tables. Pointer-log, audit, and outbox rows exist but are invisible to table semantics.
Platform operator	A hardened catalog transaction log: idempotency outcomes, pointer movement, redacted conflict proof, storage-profile and credential posture, pending outbox, replay-validation failures.
Governed agent	A narrowed access path: TypeSec decides the capability, LakeCat binds the receipt to the catalog action, Sail plans the effective scan, the agent gets bounded work instead of broad credentials.
QueryGraph importer	A proof-bearing bootstrap: table, view, management, credential, scan, commit-history, OpenLineage, and graph-import anchors that must line up before the semantic layer is accepted.

That separation also keeps standardization honest. The interesting portable shapes are never “LakeCat uses Turso” or “QueryGraph imports a bundle”; they are smaller behaviors — idempotent commit replay, catalog pointer history, governed credential vending, proof-carrying scan planning, redacted conflict evidence, lineage receipt binding. The posture, in five rules:

1. Keep base Iceberg REST behavior strict and boring.
2. Keep LakeCat proof surfaces optional for the clients that need them.
3. Keep TypeSec and QueryGraph semantics out of Iceberg metadata.
4. Use Sail for reusable table-format and planning semantics.
5. Promote only small, interoperable proof profiles once real workflows prove multiple engines and catalogs need them.

6 Why the Catalog Stays Thin

The most dangerous failure mode for a “smart” catalog is becoming a partial engine. It starts innocently — validate a schema here, expand a manifest list there, peek at a format-version field, check a delete file. Each small parser looks cheaper than an engine call. Over time the catalog grows a second Iceberg implementation with weaker tests, fewer real execution users, and quiet drift from how the planner actually behaves.




Figure 5: Diagram 5

LakeCat takes the opposite path: **the catalog owns the transaction; Sail owns the table semantics**. Sail is the right home because it is Rust-native and already holds the structures this needs — generated Iceberg REST models, table providers, manifest pruning, metadata-as-

data paths, commit plumbing, and format-version handling. Anything that needs field-id binding, schema/partition evolution, manifest metrics, delete association, row lineage, v4 metadata trees, or snapshot/branch selection moves toward Sail.

This matters because intermediation, left passive, *loses* information. A pass-through catalog sees only the table name, the pointer, the caller, and maybe a credential request; the engine sees schema, manifests, statistics, deletes, and filters; governance sees policy; lineage sees an after-the-fact event. Each system gets a shard of the truth, and the operational cost is concrete:

- policy can be checked before access but not carried *into* scan planning;
- credentials can be vended without proving why a raw-credential exception was allowed;
- lineage can say something happened but cannot bind to the exact governed plan, snapshot, policy, and metadata.

The next generation of callers — agents, notebooks, services, model pipelines — needs stronger guarantees that live precisely between catalog state and engine planning: a policy enforced before a scan is planned; column restrictions that narrow the projection before file tasks exist; policy-derived row predicates that become mandatory filters; a stateless `fetchScanTasks` that cannot widen a prior governed plan; short-lived, scoped, audited credentials; idempotent commit retries; and graph/lineage side effects that reflect committed state rather than best-effort handler work. A catalog too far from the engine can only check a policy and hand back a pointer, hoping the client preserves the restriction. LakeCat keeps the catalog boundary thin and binds its proof to Sail's plan so the restriction is real.

7 The Architecture

Thin does not mean trivial. LakeCat owns the durable catalog state that must be correct even when external sinks are down, and nothing more. Its responsibilities are exactly:

- serve the Iceberg REST Catalog API for standard clients;
- model projects, warehouses, namespaces, tables, views, and storage profiles;
- persist metadata pointers and compare-and-swap commit history;
- validate idempotency keys and replay only matching commit bodies;
- resolve request identity from headers, bearer tokens, agents, and TypeDID envelopes;
- ask TypeSec for authorization decisions and persist the receipts;
- route scan planning and commit preparation through Sail;
- record audit and outbox events inside the catalog transaction;
- drain committed events to Grust and OpenLineage sinks;
- publish a QueryGraph bootstrap bundle.

Everything else is deliberately excluded — LakeCat does not invent a table format, fork manifest pruning, own graph traversal, decide security semantics, or author QueryGraph's business model. Those belong to Sail, Grust, TypeSec, and QueryGraph respectively (Chapter 4).

7.1 The CatalogStore seam

All durable state sits behind one trait, so the memory store (for tests and embedded use) and the Turso store (for durable local deployments) are interchangeable. The interesting methods are the table lifecycle and the idempotent replay hook:

```
#[async_trait]
pub trait CatalogStore: Send + Sync + 'static {
    async fn create_namespace(&self, warehouse: &WarehouseName, namespace:
Namespace)
        -> LakeCatResult<()>;
    async fn list_tables(&self, warehouse: &WarehouseName) ->
LakeCatResult<Vec<TableRecord>>;
    async fn create_table(&self, table: TableRecord) ->
LakeCatResult<TableRecord>;
    async fn load_table(&self, ident: &TableIdent) -> LakeCatResult<TableRecord>;

    // Optimistic pointer advancement: the catalog transaction.
    async fn commit_table(&self, ident: &TableIdent, commit: TableCommit)
        -> LakeCatResult<TableRecord>;

    // Exact idempotent replay: same key + same request hash -> stored response.
    async fn replay_table_commit(
        &self,
        ident: &TableIdent,
        idempotency_key: &str,
        idempotency_request_hash: &str,
    ) -> LakeCatResult<Option<TableRecord>>;

    // Compact pointer-log history for operators and QueryGraph.
    async fn table_commit_records(&self, ident: &TableIdent, start: u64, end:
Option<u64>)
        -> LakeCatResult<Vec<TableCommitRecord>>;
}
```

7.2 The read path

A read begins like a standard catalog request and ends with a *governed* Sail plan. The policy restriction becomes part of planning, not a note beside it.




Figure 6: Diagram 6

The governing rule is **narrow, never widen**. An empty client projection under a column restriction means *the allowed columns*; a client projection may narrow further but cannot widen. LakeCat records both the requested and the effective projection (and stats fields) as replay evidence, and recomputes the restriction on every stateless `fetchScanTasks` so a stale token cannot expand back to all columns. Outbox admission enforces that the governed scan replay carries the same read-restriction at the top level and inside `authorization-receipt.context`, a nonblank purpose, and a positive `max-credential-ttl-seconds`, and rejects unknown fields — so graph, lineage, and QGLake evidence can never inherit a claim the receipt did not capture.

7.3 The commit path

The write path follows the same principle: LakeCat owns the transaction, Sail owns Iceberg validation and metadata assembly.




Figure 7: Diagram 7

Three invariant groups make this safe, each enforced once:

- **Metadata-write safety.** The new object must be a concrete child of the table's matched storage-profile prefix — never the current pointer, never the profile root, never a path with literal or percent-encoded dot segments, query strings, fragments, userinfo, or credential markers. Writes are create-only; if the store CAS loses after a write, LakeCat makes a bounded retry to delete the orphaned object and otherwise preserves the original conflict.
- **Idempotency.** Keys are 1–128 ASCII chars (Idempotency-Key / x-lakecat-idempotency-key, matching if both present). The same key + same body replays the stored response before any Sail call or write; the same key + a different body conflicts. Only audit-safe hashes are stored — never raw keys or secrets. Ordinary clients may commit without a key; they simply get no replay proof.
- **Redaction.** Every commit error — pointer conflict, overwrite, prefix mismatch, cleanup failure, backend error — reports sha256: hashes of the metadata location and failure detail, never raw object paths, storage roots, profile ids, or backend text. A conflict still looks like an ordinary Iceberg conflict to the client.

Commit records carry compact evidence — Iceberg format version, current snapshot id, request/response/idempotency/policy hashes, principal — so operators and QueryGraph can answer audit questions from the pointer log without parsing full metadata. Operators read it through a governed endpoint:

```
curl -s -H 'x-lakecat-principal: operator@example.com' \
  http://127.0.0.1:8181/management/v1/warehouses/local/namespaces/default/tables/
events/commits
```

That read itself enters the outbox as `table.commits-listed` and drains as lineage plus Commit graph anchors keyed by table and sequence — an auditable trail without giving anyone direct database access or turning LakeCat into a graph query engine.

7.4 The durable spine

The durable local store uses the Rust `turso` crate behind the `turso-local` feature. Object storage remains the source of Iceberg metadata files; Turso holds the atomic pointer and the control-plane memory: projects and warehouses, storage profiles, namespaces and tables, the metadata pointer log, idempotency records, soft deletes, policy bindings, audit events, and outbox events. This mirrors the Iceberg contract — metadata files describe the table; catalog state decides which file is current.

Two disciplines keep the spine trustworthy:

- **Decoded-row binding.** Every read reconciles the decoded JSON against the row's own key columns before returning it — a warehouse against its project and storage root, a namespace against its path, a policy binding against its table and enforced flag, a storage profile against its prefix/provider/mode, a table body against its identity. A corrupted row index can therefore never point a valid-looking record at the wrong tenant, path, or policy.
- **Strict outbox draining.** LakeCat projects a batch to graph and lineage *first*, then acknowledges the whole batch; if projection fails, nothing is acknowledged, and an under-count acknowledgement is a hard error rather than a quiet partial success. The drain refuses unknown event types (they stay pending rather than vanish), rejects event-id/payload-hash drift, and validates governed-read and commit evidence — policy-hash arrays must be non-empty full digests, top-level and receipt read-restriction must match, purpose and TTL must be present — before any sink sees the event. Malformed source evidence stays available for retry instead of being promoted into a QGLake handoff.

8 The Benchmark Suite

A catalog lives on more than one path, so LakeCat is measured on more than one. The suite covers four axes a catalog and its engine actually exercise — committing under contention, scanning cold and warm through an object-store cache, competing against a JVM engine on identical files, and carrying a *stock* Iceberg client through a full write-and-read. The commit benchmark lives in `querygraph/catalog-commit-bench`; the cache-scan, `rust-versus-jvm`, and read-write benchmarks live in the broader `querygraph/catalog-bench` suite. All of them run behind the same impartial Docker/MinIO harness, so every number compares systems doing identical work against identical storage. This chapter walks them in turn, beginning with the commit path the rest of the catalog is built to serve.

The commit path is where a catalog earns its keep, and it is exactly the part the usual benchmarks ignore. TPC-DS and TPC-H measure query engines; they touch the catalog only incidentally. To

measure the catalog itself, LakeCat is exercised by `catalog-commit-bench` — a small, catalog-agnostic driver that issues `set-properties` commits with no data files, so each request runs the commit machinery and nothing else: validate the update, write a fresh `metadata.json`, advance the metadata pointer under compare-and-swap, and persist durably. It reports two numbers — sequential per-commit latency, and concurrent throughput under eight writers contending on one table.

8.1 Measuring catalogs, not object stores

A commit-path comparison is only honest if every catalog does the same work to the same storage. The harness puts LakeCat, Apache Nessie, Apache Gravitino, and Apache Polaris on one Docker network behind one MinIO, and every catalog writes its Iceberg `metadata.json` to the same `s3://warehouse` bucket. Each catalog’s own state store — Turso for LakeCat, the version store for Nessie, the metastore for Polaris/Gravitino — is its private metadata-pointer bookkeeping, the analogue across all of them; the Iceberg metadata object itself lands in the shared object store for everyone. Without that discipline the benchmark would be comparing S3 clients, not catalogs.

8.2 Where LakeCat lands

One impartial sweep, 1000 sequential commits then eight concurrent writers for six seconds, every catalog writing to the same MinIO:

Catalog	Sequential	p50	Concurrent (8 writers)
Nessie 0.107.5	228.6 /s	4.04 ms	164.0 /s
LakeCat 0.2.0	198.2 /s	4.52 ms	311.6 /s
Gravitino	163.9 /s	5.74 ms	340.2 /s
Polaris 1.5.0	97.6 /s	9.81 ms	91.5 /s

LakeCat’s per-commit latency sits second of the four — its median is within a hair of Nessie’s and ahead of Gravitino and Polaris — and on concurrent throughput it is second only to Gravitino, well ahead of Nessie. Polaris is the heaviest per commit; its RBAC checks and credential subscoping are real governance cost, not inefficiency.

8.3 What it took to get there

The first honest LakeCat numbers were not competitive — its commit median was nearly twice the Java catalogs’. The cause was neither Rust nor the catalog’s own logic, which runs in well under a millisecond; it was two missing connection-reuse habits the JVM data ecosystem standardized decades ago. LakeCat rebuilt its S3 client — credential chain, HTTP client, a fresh connection with no keep-alive — on every commit; a request trace showed about one `PutObject` per commit at roughly 1.7 ms server-side, so most of the latency was per-commit client setup, not the write. Caching one client per bucket cut the median almost in half. Separately, the write transaction opened a new Turso connection and re-applied its MVCC pragmas on every commit; pooling pragma-warmed connections — still a distinct one per concurrent writer, so MVCC concurrency

is unchanged — cut the median again. Together these took the median commit from about 12.6 ms to about 4.5 ms, moving LakeCat from the slowest of the field to the front of it.

8.4 Audit and idempotency: the durable cost of features

The small remaining gap to Nessie is not speed; it is work LakeCat does that the others do not. Every LakeCat commit performs several writes inside one transaction: the metadata-pointer compare-and-swap, a pointer-log row, an audit event, a transactional-outbox row staged atomically with the commit — so lineage and graph events can never be lost or emitted without it — and an idempotency record, so a retried commit replays its prior result rather than double-applying. That is a durable audit trail, an atomic outbox, and idempotency, fsynced per commit — guarantees the leaner version stores do not provide in the same transaction. LakeCat is paying for the spine described in *The Architecture*, by design; the gap closes by relaxing those guarantees, not by changing languages.

This is also why Rust did not, by itself, win the benchmark. The commit path is I/O-bound, so runtime CPU speed is nearly irrelevant against a network PUT and an fsynced transaction, and a warm, long-running server running a tight commit loop is the JVM's best case — JIT-compiled hot paths and warm connection pools, with its real weaknesses of cold start and memory footprint nowhere in frame. Where the Rust implementation keeps its edge is exactly what a warm steady-state benchmark hides: no GC pauses and so steadier tail latency, a far smaller resident footprint, and instant cold start — properties that matter for serverless, edge, and many-tenant-per-host deployments rather than for a single warm server in a loop.

The driver, the impartial Docker/MinIO harness, and the full results live in `querygraph/catalog-commit-bench`.

8.5 The object-store cache and the scan benchmark

Where the commit benchmark holds the object store constant to isolate the catalog, the cache-scan benchmark does the opposite: it puts the read path under test and asks what a warm cache is worth. The answer is large, because the work a warm cache removes is a network round-trip. The benchmark plans a scan, reads it once cold straight from MinIO, then reads it again with Sail's object-store cache warm, over the same 87 MB dataset.

The cache is a per-worker, read-through *page* cache in Sail's `sail-object-store` crate — `CachingObjectStore` over a `CacheConfig` — added to answer Sail issue #1015. It is ported from `lancedb/ocra`, attributed in the crate, with the original Moka backing store swapped for Foyer. It is opt-in: `SAIL_OBJECT_STORE_CACHE` turns it on, and `SAIL_OBJECT_STORE_CACHE_PAGE_SIZE`, `_MEMORY`, and `_METADATA` tune it — by default 1 MiB pages, 1 GiB of value memory, and 64 MiB of metadata. Because `object_store` exposes its read methods as a non-overridable blanket trait, the cache cannot wrap them directly; it intercepts the two range entry points the engine actually reads through — `get_opts` and `get_ranges` — and serves whole pages from memory. The current tier is in-memory only; Foyer's hybrid disk tiering is a follow-up the same seam already anticipates.

On the same files the difference is roughly twenty-six fold: the per-file scan median falls from about 47.5 ms cold and uncached to 1.81 ms warm. None of this lives in LakeCat. A read-through

object-store cache is a reusable engine concern, so it belongs in Sail; LakeCat consumes it through the same dependency seam it uses for planning and benefits without owning a line of cache code. How that seam is wired — and which Sail commits carry the cache — is the subject of LAKECAT-SAIL.md.

8.6 Rust versus the JVM

The commit chapter was careful about what Rust did and did not buy: on an I/O-bound, fsynced commit loop against a warm JVM, runtime CPU speed barely registers. The rust-versus-jvm benchmark asks the same question honestly on the read path, where a vectorized engine has more room to move. It runs Sail/DataFusion and Spark 3.5.3 over the identical query and the identical files.

The honest engine edge is modest: about 1.63× — a 446 ms Rust median against 729 ms for warm Spark. That is the real, like-for-like number, and it is the one to quote, because the JVM here is in its best case: warm, JIT-compiled, hot pools, doing exactly the same scan. The figure that looks spectacular — about 57.5× — is the *cached* comparison, and it earns an asterisk: the warm Foyer cache is doing most of that work, not the language. The lesson is the commit benchmark's, mirrored: the large win comes from a system-design choice — a cache living in the engine — not from Rust itself. What Rust keeps independently is what a warm steady-state benchmark hides — steadier tail latency with no GC pauses, a smaller resident footprint, and instant cold start.

8.7 The read-write benchmark and the stock-client round-trip

The most demanding test in the suite is also the simplest to state: can an *unmodified* Iceberg client create a table, write real data, and read it back through LakeCat? The read-write benchmark drives stock PyIceberg with no LakeCat-aware shim — a plain RestCatalog pointed at the service. When the suite first ran this probe the answer was no; an early phase had recorded the full stock round-trip as impossible. The value of the benchmark was that it named *why*, gap by gap, and each gap turned out to be small and closable.

Five fixes, four in LakeCat and one in Sail, turned impossible into routine:

- **Map fields became objects.** LakeCat had been serializing Iceberg's map-typed fields — defaults, overrides, config, properties — as arrays of {key, value} pairs, which a stock client cannot parse. A config_map serde adapter now emits them as JSON objects, so a plain RestCatalog reads /config and proceeds.
- **/config advertises canonical routes.** The endpoint now lists spec-canonical <METHOD> /v1/{prefix}/... forms — keeping the legacy strings alongside — so a stock client routes its calls where the spec says they live.
- **listTables exists.** A GET .../namespaces/{namespace}/tables endpoint was added, which stock clients call as a matter of course.
- **The default build refuses what it cannot keep.** Without sail-local, a commit can validate but not truly apply table-metadata updates; it used to accept them and silently drop the change. It now *rejects* updates it cannot apply, the same fail-closed posture the deferred scan seam already takes. Real persistence is the sail-local build, which applies updates through Sail's apply_table_updates.

- **Sail learned the append.** `apply_table_updates` now handles `add-snapshot` and `set-snapshot-ref`, the two updates a data append produces — so a real snapshot commit lands as new table metadata rather than being refused.

The split is the boundary doing its job: compatibility shape — object-typed maps, canonical routes, the missing endpoint, the honest default-build refusal — is LakeCat’s to fix, while table-format evolution, applying a snapshot append, is Sail’s. With all five in place the round-trip is real and reproducible: stock PyIceberg 0.11.1, no shim, initializes a `RestCatalog`, creates a table, calls `table.append` and gets a genuine snapshot (`snapshots_after = 1`), then scans 1000 rows back, all against a `sail-local` LakeCat. With the Foyer cache warm the read side of that loop runs about 150× faster than cold. What had been declared impossible is now a benchmark that passes on every run.

The cache-scan, rust-versus-jvm, and read-write drivers, the MinIO harness, and the stock-client round-trip live in `querygraph/catalog-bench`; the integration seam they exercise — how LakeCat consumes Sail, and which Sail commits the round-trip depends on — is documented in `LAKECAT-SAIL.md`.

9 The Siblings and the Engine Path

Chapter 4 named the four sibling repositories. This chapter shows how LakeCat talks to each — what it hands off, and what it deliberately refuses to own.

9.1 Grust: the graph boundary

Catalog events naturally form a graph: a server contains projects, a project contains warehouses, a warehouse contains namespaces and credential-rooted storage profiles, a namespace contains tables, and a table has columns, snapshots, commits, policies, scan plans, principals, and lineage runs. `QueryGraph` wants that graph — but LakeCat must not become a graph database.

LakeCat emits a bounded envelope of stable catalog facts through a catalog-facing sink; Grust owns the taxonomy, stores, traversal, and Cypher:

Server CONTAINS Project	Warehouse HAS_STORAGE_PROFILE StorageProfile
Project CONTAINS Warehouse	Table GOVERNED_BY Policy
Warehouse CONTAINS Namespace	Principal CAN_PLAN ScanPlan
Namespace CONTAINS Table	Commit DERIVED_FROM Snapshot
Table HAS_COLUMN Column	LineageRun USED_BY QueryGraphModel

High-cardinality file and manifest facts stay queryable through Sail’s metadata-as-data rather than being smuggled into the event sink. Storage-profile and credential-vend events replay with redacted evidence only — `secret-ref-present`, the provider, never the secret URI — so `QueryGraph` can see *that* a principal attempted credential-root access without seeing the credential. The `grust-turso-local` feature proves the boundary end to end: LakeCat writes catalog events into a Grust-owned Turso graph store and Grust Cypher reads them back, with LakeCat never parsing Cypher or executing traversal. Ocelot consumes the published **Grust 0.12.0, Lobster** crates for both the graph API and Turso backend; the feature name describes LakeCat’s in-process integration row, not a path checkout.

9.2 TypeSec: the authorization boundary

LakeCat is a policy *enforcement* point, not the author of security semantics. Every externally meaningful action — config reads; namespace and table lifecycle; scan planning; commit; credential vending; policy management; graph and lineage reads — passes through TypeSec. LakeCat gathers context (principal and agent DID, bearer subject, warehouse/namespace/table, columns, snapshot, requested credential duration, purpose, active bindings), asks TypeSec for a decision, persists the receipt with audit-safe hashes, and applies the restriction *before* Sail plans.

This is where ODRL becomes operational. A policy may say a principal reads only certain columns, only for a purpose, or only under a row predicate. LakeCat parses the minimal enforceable subset — accepting camel-case, kebab-case, and JSON-LD operand forms — and **fails closed**: missing or deny-shaped operators, blank allowed-column lists, blank purposes, or disagreeing purpose sources are rejected rather than guessed. Composition and reasoning stay in TypeSec; LakeCat does not grow a parallel security language.

Credential vending is the audited exception, not the default. Governed Sail-planned reads are the norm. When credentials must issue, TypeSec checks the `credentials.issue` capability for the exact secret reference and LakeCat returns only scoped, short-lived configuration — capping TTL to the tightest `max-credential-ttl-seconds` across all policy locations, replacing any issuer-supplied `lakecat.*` evidence with catalog-derived values, and recording only hashed prefixes in audit. The replay-admission rules from Chapter 5 (matching top-level and receipt restrictions, full digests, closed field sets) apply identically here, so a credential or raw-exception event cannot drift before it reaches graph, lineage, or QGLake proof.

Ocelot consumes published **TypeSec 0.12.0, Torcello**. LakeCat owns the enforcement point and durable receipt binding; TypeSec owns authorization, capability, and receipt semantics.

TypeSec now also owns the completed Marciana v1 security boundary for governed agent memory. A `MemoryVault` requires operation-specific capabilities for remember, recall, and forget; applies purpose and clearance policy before rehydrating content; and preserves labels, quarantine, temporal history, and audit. Grust’s `querygraph-memory` adapter owns the opaque record and entity graph on durable Turso/libSQL storage. At the service edge, `qg-rust` verifies the signed `TypeDID` envelope and uses its signing `did:key`—not a caller-supplied body field—as the TypeSec policy subject.

That memory is not LakeCat’s `MemoryCatalogStore`. The latter is an operational implementation of `CatalogStore` for tests and embedded use, interchangeable with LakeCat’s durable Turso catalog store. It holds namespaces, tables, views, policies, receipts, idempotency state, and outbox events; it is not an agent recall system. LakeCat contributes governed catalog and lineage evidence that a `QueryGraph` application may later choose to remember. It does not mint Marciana capabilities, store Marciana records, or reveal recalled content.

9.3 Sail and the v3→v4 path

Sail is a Rust-native engine — Arrow, DataFusion, generated Iceberg REST models, catalog-provider seams, manifest pruning, metadata-as-data — so LakeCat can *ask* for typed Iceberg

behavior instead of parsing just enough JSON to survive. The questions LakeCat sends to Sail are the ones that require table-format knowledge:

- which field ids satisfy this projection?
- which filters are enforceable at planning time?
- which manifests and files survive pruning?
- which delete files must accompany the selected data files?
- which scan tasks are children of a governed parent plan?
- which v4 fields are known, preserved as passthrough, or not yet safe to read?

For Ocelot, the Sail crates report version 0.6.4 and resolve from the lakecat branch of github.com/querygraph/sail. Cargo.lock pins that branch to revision bddb1706ba2308e5029d47f04f03121236edbfa6. The branch is explicit release provenance: it carries the Iceberg planning, snapshot-append, and Foyer cache work LakeCat needs until those APIs can move to an upstream or published Sail line.

LakeCat evolves under three rules: conform to Iceberg v3 for ordinary clients; preserve unknown/emerging v4 metadata without claiming settled semantics; prefer typed Sail support the moment it exists, using JSON passthrough only as a bridge. Today the v4 bridge is deliberately narrow — when LakeCat sees `format-version: 4`, lakecat-sail extracts the stable envelope (table UUID, location, schema id, snapshot id, sequence number, manifest-list path, default spec, field names), can plan a governed manifest-list scan task from it, and re-validates the signed plan task on `fetchScanTasks` so a stateless fetch can neither drift to a different manifest list nor widen the governed projection. Pruning and typed metadata-tree semantics wait for Sail-owned v4 support.

The handoff between LakeCat and Sail should therefore be compact and typed:

Work item	LakeCat responsibility	Sail responsibility	Durable proof Lake-Cat should keep
Table load	Resolve warehouse, namespace, table, principal, and current metadata pointer.	Interpret table metadata and expose table status using reusable Iceberg types.	Table identity, pointer hash, format version, principal, and receipt/action context.
Governed scan	Ask TypeSec for the effective restriction and bind it to the request.	Bind projection and predicates to field ids, prune manifests and files, attach deletes, and produce scan tasks.	Requested and effective projection hashes, predicate hash, snapshot, task count, delete posture, plan hash, and receipt hash.
Fetch scan task	Check that a stateless task fetch belongs to the prior governed plan.	Reapply task interpretation without widening the plan.	Prior plan hash, task token hash, effective stats-field evidence, and receipt hash.
Commit	Own CAS, idempotency, pointer-log, audit, and outbox state.	Validate table metadata, commit requirements, format-version behavior, and future v4 typed semantics.	Expected and new pointer hashes, format version, snapshot evidence, idempotency hash, audit id, and outbox id.
Metadata-as-data	Authorize the request and record who asked.	Expose snapshots, manifests, files, deletes, partitions, and history as engine-shaped metadata views.	Metadata view name, pointer hash, result-shape hash, principal, and receipt hash.

9.4 The semantic handoff: Croissant, CDIF, OSI, OpenLineage

QueryGraph needs a semantic picture, not just physical access. LakeCat publishes one without pretending to be QueryGraph. The bootstrap bundle carries Semantic Croissant and CDIF projections (dataset/field discovery), an OSI handoff of stable anchors, ODRL artifacts and TypeSec policy context, OpenLineage events for catalog changes and plans, a Grust-ready graph envelope, and a manifest that hashes every artifact. The boundary is careful: LakeCat publishes stable anchors and governed source metadata; QueryGraph authors the business model. Because OpenLineage events drain from the durable outbox in `created_at,event_id` order *after* the catalog transaction, lineage reflects committed state rather than a handler’s best-effort side effect.

9.5 The full stack

When LakeCat is done, a standard engine still loads and commits tables without knowing Query-Graph exists, while governed callers get the richer path:




Figure 8: Diagram 8

9.6 Ocelot's dependency baseline

The release uses one coordinated substrate wave. The manifest expresses desired versions and sources; the lockfile records the exact build:

Component	Ocelot baseline	Relationship to LakeCat
LakeCat workspace	0.3.0, Ocelot	All LakeCat crates and qglake-bundle share the workspace version.
Grust	0.12.0, Lobster, crates.io	grust-graph and grust-turso provide graph projection, Cypher, and the durable Turso graph backend.
TypeSec	0.12.0, Torcello, crates.io	The typesec facade provides the authorization and receipt boundary used by lakecat-security.
Sail	0.6.4, git branch querygraph/sail#lakecat, locked at bddb1706	Sail supplies Iceberg planning, model, commit-update, and object-store-cache behavior not yet consumed from a published release.
QueryGraph	0.4.0, Sentinel	Acceptance peer rather than a compiled LakeCat dependency; its importer consumes LakeCat qglake-bundle and lakecat-core 0.3.0 while the handoff gate drives verify/import under --locked.

This distinction matters operationally. Updating a version in prose is not a dependency upgrade. Ocelot's reproducible statement is the combination of `Cargo.toml`, `Cargo.lock`, the local dependency contract, and the locked QueryGraph handoff.

9.7 Implementation shape

The workspace expresses the architecture directly:

```
crates/
lakecat-core      stable IDs, errors, time, config, content hashes
lakecat-api       Iceberg REST request/response adapters
lakecat-store     catalog state traits + Turso-backed implementation
lakecat-sail      Sail provider bridge and privileged planning client
lakecat-graph     catalog-facing Grust sink/adapters
lakecat-security  TypeSec integration and authorization receipts
```

lakecat-lineage	OpenLineage projection and event receipts
lakecat-querygraph	Croissant/CDIF/OSI/ODRL/OpenLineage bootstrap projection
qglake-bundle	portable bootstrap wire types, hashes, and validators
lakecat-service	axum service, middleware, auth, routing
lakecat-cli	admin, local demo, conformance, bootstrap export

Feature gates keep integrations honest, and embedded defaults stay safe for tests so a memory-store test never accidentally depends on a sibling repo:

sail-local	in-process Sail APIs for planning and provider integration
typesec-local	in-process TypeSec APIs for governance and TypeDID
verification	
grust-local	in-process Grust APIs for catalog graph projection
grust-turso-local	Grust's Turso backend for durable catalog graph projection
turso-local	the Turso-backed durable store

The runtime honors the same line: without `sail-local`, the deferred Sail seam *rejects* scan planning rather than fabricating an empty plan, so any real read reflects the engine that interprets Iceberg metadata, never a catalog-shaped placeholder. Standard compatibility lives at `/catalog/v1`; management APIs, `/querygraph/v1/bootstrap`, and feature-gated Sail planning sit beside it, never in front of it. (The Ocelot-era release gate and dependency contract that hold this together are covered in the release chapter.)

10 Workflow Examples

The catalog is easiest to understand by watching it participate in ordinary work. LakeCat should not ask users to think about graph, lineage, security, and Sail every time they read a table. Those systems should appear when they matter: at the boundary where a name is resolved, a policy is enforced, a plan is created, credentials are withheld or issued, and a durable event is replayed.

The examples below use one table, `local.default.events`, but the pattern is the same for larger warehouses. The important point is not the exact sample data. It is the catalog role in each workflow. The dense replay and handoff checks are in the appendices so the examples can stay human-readable.

10.1 Starting The Catalog

A local operator starts LakeCat as an Iceberg REST catalog plus management surface:

```
cargo run -p lakecat-service --features sail-local,turso-local,typesec-local,grust-local
```

The standard catalog path is still `/catalog/v1`. The management and QueryGraph surfaces sit beside it:

```
/catalog/v1
/management/v1
/querygraph/v1/bootstrap
```

A simple configuration read shows the split. Standard engines care about the Iceberg endpoints. Operators and QueryGraph care about management and bootstrap surfaces. The service binds

127.0.0.1:8181 by default; set LAKECAT_BIND_ADDR to change the address (every example below assumes the default).

```
curl -s http://127.0.0.1:8181/catalog/v1/config
```

The defaults intentionally separate compatibility from future capability:

```
{
  "defaults": {
    "lakecat.compatibility": "iceberg-rest",
    "lakecat.format.baseline": "iceberg-v1-v3",
    "lakecat.format.v4": "extension-ready",
    "lakecat.format.v4.bridge": "json-passthrough",
    "lakecat.format.v4.typed-sail": "unavailable"
  }
}
```

The Iceberg REST map-typed fields (defaults, overrides, table/credential config, and namespace properties) are JSON objects on the wire so that stock clients (pyiceberg/Spark/Trino) parse them directly; the advertised endpoints include the spec-canonical <METHOD> /v1/{prefix}/... forms alongside LakeCat's legacy strings.

That means LakeCat can preserve and replay emerging v4 metadata through the Sail JSON bridge, but it is not claiming typed Sail v4 semantics yet. The same posture is captured in replay evidence: config defaults and overrides stay ordinary duplicate-free key/value entries, advertised endpoints prove that the standard REST and additive governed surfaces were present, and QueryGraph sees compact config proof only after service replay accepts it.

10.2 Registering The Warehouse Shape

An operator usually starts with management objects. A server groups projects. A project groups warehouses. A warehouse owns namespaces, tables, views, storage profiles, policy bindings, and the metadata pointer state that standard engines see through Iceberg REST.

```
curl -s -X PUT http://127.0.0.1:8181/management/v1/servers/prod \
-H 'content-type: application/json' \
-d '{
  "display-name": "Production LakeCat",
  "endpoint-url": "https://lakecat.example.com",
  "properties": {"owner": "platform"}
}'
```

```
curl -s -X PUT http://127.0.0.1:8181/management/v1/projects/resilience \
-H 'content-type: application/json' \
-d '{
  "display-name": "Resilience Desk",
  "server-id": "prod",
  "properties": {"environment": "demo"}
}'
```

```
curl -s -X PUT http://127.0.0.1:8181/management/v1/projects/resilience/
```

```

warehouses/local \
-H 'content-type: application/json' \
-d '{
  "display-name": "Local QGLake Warehouse",
  "storage-root": "file:///tmp/lakecat/qglake",
  "properties": {"querygraph": "enabled"}
}'

```

These writes are not Iceberg table metadata. They are catalog control-plane state. LakeCat persists them durably, records authorization receipts, and writes outbox events. When the outbox drains, server, project, warehouse, storage-profile, and policy-binding changes become catalog graph events and OpenLineage receipts. QueryGraph can later learn the management shape without requiring every Iceberg client to understand it.

The replay rule is simple: management evidence is closed, attributable, scoped, and redacted. Tenant-root records bind decoded IDs back to their route scope, policy-binding replay carries the ODRL content hash, namespace/table/view lifecycle replay carries the matching action receipt, and raw endpoint or storage-root material is replaced by full hash evidence before graph, OpenLineage, QGLake, or QueryGraph sees it.

10.3 Storage Profiles And Credential Roots

Storage profiles bind a warehouse to physical storage roots and credential issuance policy. A local profile can return scoped local file configuration. A remote profile should usually reference a secret store and require TypeSec to authorize issuance before any resolver sees the secret reference.

```

curl -s -X PUT \
  http://127.0.0.1:8181/management/v1/warehouses/local/storage-profiles/local-
events \
-H 'content-type: application/json' \
-d '{
  "location-prefix": "file:///tmp/lakecat/qglake/events",
  "provider": "file",
  "issuance-mode": "local-file-no-secret",
  "public-config": {"lakecat.purpose": "developer-loop"}
}'

```

```

curl -s -X PUT \
  http://127.0.0.1:8181/management/v1/warehouses/local/storage-profiles/s3-events
\
-H 'content-type: application/json' \
-d '{
  "location-prefix": "s3://lakecat/events",
  "provider": "s3",
  "issuance-mode": "short-lived-secret-ref",
  "secret-ref": "vault://kv/lakecat/events",
  "public-config": {
    "lakecat.region": "us-west-2",

```

```

    "lakecat.purpose": "production-events"
  }
}'

```

The store accepts only plainly addressed roots: no query strings, fragments, URI userinfo, literal or percent-encoded dot segments, provider/location mismatches, ambiguous longest-prefix ties, secret-looking public-config, or decorated secret-ref values. Operator-facing errors and replay evidence use hashes such as storage-profile-prefix-hash, location-prefix-hash, secret-ref-hash, and public-config-key-hash rather than echoing submitted paths or secret-like strings.

The catalog row stores the public profile and secret reference, not raw cloud keys. A later credential request is checked against TypeSec and against the effective read restriction for the target table. Agents with fine-grained table restrictions are steered to governed Sail-planned reads instead of raw credentials. Trusted humans can receive audited standard credentials only when policy allows the exception.

For production-shaped secret managers, LakeCat keeps a provider-dispatch seam instead of embedding credentials in catalog state. `vault://` can use the built-in Vault HTTP backend. `aws-sm://`, `gcp-sm://`, and `azure-kv://` can dispatch to configured provider backends after TypeSec authorizes the exact secret-ref resource, and local file-backed roots let the same dispatch shape run in a single-node loop. Returned credentials must stay inside the selected storage profile prefix, carry the policy-derived TTL cap, and preserve LakeCat-owned profile/provider/mode/principal evidence. Replay records redacted prefix and issuer hashes, not credential material.

10.4 A PySpark User Reads Iceberg

A PySpark user should not need to know about QueryGraph. They configure Spark's Iceberg REST catalog and point it at LakeCat:

```

from pyspark.sql import SparkSession

spark = (
    SparkSession.builder
        .appName("lakecat-events")
        .config("spark.sql.catalog.lakecat", "org.apache.iceberg.spark.SparkCatalog")
        .config("spark.sql.catalog.lakecat.type", "rest")
        .config("spark.sql.catalog.lakecat.uri", "http://127.0.0.1:8181/catalog/v1")
        .config("spark.sql.defaultCatalog", "lakecat")
        .getOrCreate()
)

events = spark.table("default.events")
events.select("event_id", "severity").where("severity = 'critical']").show()

```

For an unrestricted principal, the flow looks like a normal Iceberg read:

1. Spark asks LakeCat to resolve `default.events`.
2. LakeCat checks the principal and table capability.

3. LakeCat returns an Iceberg-compatible table response.
4. Spark plans the read using Iceberg metadata.

For a governed principal, the more interesting path is server-side planning. The user request still looks ordinary, but LakeCat derives the mandatory restriction before Sail sees the plan:

```
client asks:    event_id, severity, raw_payload
policy allows:  event_id, severity
Sail receives:  event_id, severity
```

The catalog does not trust the client to remember that. The restriction is re-applied when scan tasks are fetched, and the audit payload records the policy hash, narrowed columns, row predicate, storage location, metadata location, principal, requested stats fields, and effective stats fields. QGLake handoff proof compares planned and fetched evidence so a saved artifact cannot widen the server-derived purpose, TTL cap, allowed columns, row predicate, stats fields, or policy hashes.

10.5 A Notebook Requests Credentials

Credential vending is deliberately different from scan planning. Returning storage credentials gives the client broader power than returning a governed task list, so LakeCat treats it as an exception path.

```
curl -s \
  -H 'x-lakecat-principal: agent:triage' \
  http://127.0.0.1:8181/catalog/v1/local/namespaces/default/tables/events/
credentials
```

For an agent bound by a fine-grained restriction, LakeCat should fail closed:

```
{
  "credentials": [],
  "lakecat:credential-block-reason": "fine-grained read restriction requires
Sail-planned reads"
}
```

That empty credential response is not a missing feature. It is the intended agentic posture. The agent should ask LakeCat to plan the read through Sail, not receive raw storage reach. For a trusted human principal, policy can allow an audited raw credential exception, but the audit and replay evidence still carry the principal, table, decision, reason, TTL cap, storage-profile anchor, and redacted returned-prefix hashes.

10.6 A View Becomes Part Of The Catalog Story

Views are catalog objects too. LakeCat stores durable view records with SQL, dialect, schema version, typed columns, properties, creator, and warehouse scope. They can be managed through the management API or through warehouse-prefixed catalog routes.

```
curl -s -X PUT \
  http://127.0.0.1:8181/management/v1/warehouses/local/namespaces/default/views/
events_view \
```

```

-H 'content-type: application/json' \
-d '{
  "sql": "select event_id, severity from default.events where severity =
'\''critical'\''",
  "dialect": "spark-sql",
  "schema-version": 1,
  "columns": [
    {"name": "event_id", "data-type": {"type": "long"}, "nullable": false},
    {"name": "severity", "data-type": {"type": "string"}, "nullable": false}
  ],
  "properties": {"owner": "resilience-desk"}
}'

```

This is not Iceberg business metadata glued onto a table. It is catalog state about a view object. LakeCat records view.listed, view.upserted, view.loaded, and view.dropped audit events. Outbox replay projects listing reads to namespace-scoped graph and OpenLineage evidence, and projects single-view changes and reads to catalog-facing View graph events plus LakeCat OpenLineage view dataset receipts. QueryGraph bootstrap can then include views with OSI hashes, store-assigned view versions, view-aware graph edges, and OpenLineage view counts.

A caller that wants optimistic view behavior includes the version it believes is current:

```

curl -s -X PUT \
  http://127.0.0.1:8181/management/v1/warehouses/local/namespaces/default/views/
events_view \
-H 'content-type: application/json' \
-d '{
  "sql": "select event_id from default.events where severity =
'\''critical'\''",
  "dialect": "spark-sql",
  "schema-version": 2,
  "expected-view-version": 1
}'

```

If another writer has already advanced the view, LakeCat returns a conflict before it replaces the current view or appends a receipt. Drop can use the same guard:

```

curl -s -X DELETE \
  'http://127.0.0.1:8181/management/v1/warehouses/local/namespaces/default/views/
events_view?expected-view-version=2'

```

LakeCat also writes a compact view-version receipt in the durable store. The receipt records the stable view id, assigned version, previous version, previous receipt hash, content hash, principal, operation, and timestamp. A create, drop, and recreate sequence therefore looks like version 1 upsert, version 1 drop tombstone, version 2 upsert linked to the tombstone receipt, not two unrelated version-1 chains for the same stable view id.

QueryGraph and operators can read the compact receipt chain directly from the governed management surface:

```
curl -s \
  http://127.0.0.1:8181/management/v1/warehouses/local/namespaces/default/views/
events_view/version-receipts \
  -H 'x-lakecat-agent-did: did:example:resilience-agent'
```

The response is catalog evidence, not Iceberg table metadata. It lets QueryGraph verify the version chain, including tombstones after the current view row is gone, while keeping the richer view history model available for a future Sail-owned implementation. Appendix C captures the exact receipt-chain rules.

10.7 QueryGraph Bootstrap

QueryGraph should import LakeCat facts through a verified handoff, not by scraping service internals. The bootstrap endpoint publishes a bundle with catalog tables, views, policy bindings, graph artifacts, OpenLineage artifacts, Croissant/CDIF/OSI/ODRL projections, and a manifest that hashes what was emitted:

```
curl -s \
  -H 'x-lakecat-principal: agent:querygraph-importer' \
  http://127.0.0.1:8181/querygraph/v1/bootstrap \
  -o target/qglake/lakecat-bootstrap.json
```

The exported graph includes the tenant spine:

```
Catalog HAS_SERVER Server
Server HAS_PROJECT Project
Project HAS_WAREHOUSE Warehouse
Warehouse HAS_NAMESPACE Namespace
Namespace HAS_TABLE Table
```

When LakeCat has durable management rows, those graph nodes come from the stored ServerRecord, ProjectRecord, and WarehouseRecord. Replay evidence redacts storage roots and server endpoint URLs into hashes, so QueryGraph can prove which management state was observed without inheriting local paths, bucket roots, query tokens, URI fragments, or userinfo. The tenant anchors and warehouse-to-namespace edges are manifest-covered, so an importer can reject a bundle whose namespace is detached from the warehouse or rebound to a different durable tenant chain.

The QueryGraph side verifies the bundle before importing it:

```
cd /Users/alexysrc/querygraph/qg-rust

cargo run -- lakecat-verify \
  --bundle /Users/alexysrc/lakecat/target/qglake/lakecat-bootstrap.json

cargo run -- lakecat-import \
  --bundle /Users/alexysrc/lakecat/target/qglake/lakecat-bootstrap.json \
  --output .querygraph/lakecat/import-plan.json
```

The importer checks the outer bundle hash, manifest hashes, QueryGraph-import compatibility hash, graph hash, OpenLineage hash, view receipt evidence, and the accepted receipt-chain hash

for each exported view. Graph validation belongs on the QueryGraph/Grust side; LakeCat is responsible for producing the clean catalog-facing graph projection and the replayable anchors.

For the full local handoff, LakeCat carries a script that runs both sides without writing generated artifacts into the QueryGraph checkout:

```
scripts/qglake-handoff-local.sh
```

The script starts LakeCat on 127.0.0.1:18181, uses a Turso-backed local store, generates the paired QGLake bootstrap bundle and lineage-drain response, runs `qglake-verify-replay`, then runs QueryGraph's `lakecat-verify` and `lakecat-import` against the same bundle. The compact `handoff-summary.json` records the catalog URL, principal, table scope, LakeCat replay status, QueryGraph-verified table/view counts, and semantic bundle/graph/OpenLineage/import hashes plus accepted standards. Appendix D is the detailed handoff proof map.

10.8 Draining The Outbox

LakeCat records side effects as durable outbox events. Draining the outbox is what turns committed catalog facts into graph and lineage receipts:

```
curl -s -X POST \
  -H 'x-lakecat-principal: agent:lineage-drainer' \
  http://127.0.0.1:18181/management/v1/lineage/drain \
  -o target/qglake/lineage-drain.json
```

A useful drain response includes delivered event types, graph projection counts, lineage projection counts, receipt hashes, and the authorization proof for the drain request itself. Reading the replay stream is privileged, so LakeCat records that the drainer was allowed to read lineage evidence.

The end-to-end result is a chain:

```
catalog write
-> audit event
-> outbox event
-> graph projection
-> OpenLineage projection
-> QueryGraph import evidence
```

If graph or lineage sinks are down, catalog state should not be lost or rolled back accidentally. The outbox lets LakeCat retry projection from committed state. A drain acknowledges delivery only after every projection in the batch succeeds. Replay order is part of that contract: LakeCat selects undelivered outbox events by `created_at`, `event_id` in both embedded memory tests and the durable Turso store, and duplicate or malformed pending records fail with hash-only evidence before graph emission, lineage emission, or acknowledgement.

10.9 An Agentic QGLake Flow

The agentic path is the reason LakeCat has to be more than a passive catalog. Imagine a resilience supervisor agent investigating incidents:

1. The supervisor delegates table triage to a specialist agent.
2. The specialist asks LakeCat to plan a scan over `local.default.events`.
3. LakeCat resolves the agent identity and TypeDID context.
4. TypeSec authorizes the table scan and returns a restricted capability.
5. LakeCat narrows the projection and appends the required row predicate.
6. Sail plans against the current Iceberg metadata and delete manifests.
7. LakeCat returns governed plan and fetch-task responses.
8. The specialist summarizes only the allowed result shape.
9. LakeCat records scan and credential decisions into audit/outbox.
10. QueryGraph imports graph, policy, lineage, and bootstrap evidence.

Marciana permits a separate, optional continuation after that governed result:

11. An authorized specialist sends only the allowed summary to `qg-rust`'s signed memory endpoint; TypeSec mints the write capability and the Grust adapter persists the record.
12. `qg-rust` stops and restarts while the Turso-backed memory survives.
13. A differently credentialed supervisor recalls the record because its verified DID has the required read capability.
14. A validly signed but unassigned outsider receives a TypeSec 403 denial; an unsigned caller fails envelope authentication with 401.

The key point is the absence of raw storage reach. The specialist agent does not need broad cloud credentials to do its job. It needs a governed plan, a bounded task set, and a receipt trail.

The existing QGLake fixture ends at the verified handoff in step 10. It does not run the Marciana restart proof, and its successful output must not be cited as memory verification. The separate `qg-rust/qg-python` demonstration proves steps 11–14. LakeCat supplies the governed result and its evidence anchors; `qg-rust`, TypeSec, and Grust own identity binding, recall authorization, and durable agent memory respectively.

The local fixture compresses this story into a short artifact-producing sequence:

```
cargo run -p lakecat-cli --features qglake-fixture -- qglake-fixture \
  --output target/qglake/lakecat-bootstrap.json \
  --drain-output target/qglake/lineage-drain.json \
  --principal did:example:agent
cargo run -p lakecat-cli -- qglake-verify-replay \
  --bundle target/qglake/lakecat-bootstrap.json \
  --drain target/qglake/lineage-drain.json \
  --principal did:example:agent
scripts/qglake-handoff-local.sh
cargo run -p lakecat-cli -- qglake-verify-handoff \
  --summary target/qglake-handoff/handoff-summary.json \
  --json
```

The fixture creates the sample table shape, installs a restricted policy, verifies governed scan planning, verifies fetch-scan-task reapplication, exercises delete manifest handling, probes credential-vend behavior for agents and trusted humans, verifies compact table commit-history

evidence, exports QueryGraph bootstrap artifacts, drains the outbox, and proves the resulting bundle through QueryGraph's Rust verifier/importer. It is small, but it is not decorative. It is the acceptance story for a catalog that participates in the user workflow from notebook to agent.

11 Operating The Book's Example System

The local development posture is intentionally small:

```
cargo run -p lakecat-cli -- config
cargo run -p lakecat-cli -- storage-profile-list
cargo run -p lakecat-cli -- policy-list
cargo run -p lakecat-cli --features qglake-fixture -- qglake-fixture \
  --output target/qglake/lakecat-bootstrap.json \
  --drain-output target/qglake/lineage-drain.json \
  --principal did:example:agent
cargo run -p lakecat-cli -- qglake-verify-replay \
  --bundle target/qglake/lakecat-bootstrap.json \
  --drain target/qglake/lineage-drain.json \
  --principal did:example:agent
scripts/qglake-handoff-local.sh
cargo run -p lakecat-cli -- qglake-verify-handoff \
  --summary target/qglake-handoff/handoff-summary.json \
  --json
cargo run -p lakecat-cli -- bootstrap-export \
  --output lakecat-bootstrap.json
```

The important thing is what these commands exercise. They are not a separate product surface. They touch the same catalog, policy, bootstrap, and QueryGraph export contracts that the service uses.

12 Standards And Engine Boundary Decision Record

This chapter is the release decision record for the current catalog concepts. It answers three questions that come up whenever LakeCat vocabulary gets close to Iceberg vocabulary:

1. Which terms are standard Iceberg parlance?
2. Which terms are LakeCat, QueryGraph, TypeSec, Grust, or Sail extensions?
3. Which ideas are plausible future Iceberg-adjacent optional profiles?

The distinction matters because LakeCat should be ambitious without making ordinary Iceberg clients pay for that ambition. A PySpark job should be able to load and commit a table through normal REST catalog behavior. A governed agent should be able to receive a TypeSec-backed, Sail-planned, replayable proof of bounded work. QueryGraph should be able to import QGLake evidence, OpenLineage anchors, view receipt chains, credential posture, management proof, and commit proof. Those are three different audiences using one catalog foundation.

The implementation can be pictured as four concentric contracts:

```
standard Iceberg client contract
  namespace, table, metadata location, optimistic commit
```

LakeCat catalog-control contract

identity, tenancy, Turso rows, CAS, idempotency, audit, outbox, replay

TypeSec and Sail governed-work contract

capability decision, restriction, receipt, projection, predicate, scan plan

QueryGraph semantic contract

QGLake handoff, OpenLineage, Croissant, CDIF, OSI, ODRL, graph, agents

The outer layers are additive. They can observe, prove, and compose catalog state, but they must not redefine the inner standard table contract.

12.1 Catalog Concepts In Plain Terms

This section names the catalog concepts directly, because the words can sound similar even when they live at different layers. The simplest way to read the architecture is this:

- Iceberg defines the portable table contract.
- LakeCat defines the durable catalog-control contract around that table.
- TypeSec defines the policy and capability evidence used by governed work.
- Sail defines the engine-grade interpretation of Iceberg table state.
- QueryGraph defines the semantic workflow built from accepted catalog proof.

The Rust service/catalog spine is a LakeCat implementation choice, not an Iceberg concept. Iceberg does not care whether a catalog is written in Rust, Java, Go, or Python. Iceberg cares that a client can ask for catalog config, list namespaces, load a table, and commit a new metadata pointer under ordinary REST catalog rules. LakeCat chooses Rust because the proof path should be short, typed, and close to the engine. In one process, LakeCat can parse the request, resolve identity, check tenancy, call TypeSec-style authorization, ask Sail for table interpretation, update the `CatalogStore`, append audit and outbox rows, and return the standard response. The portable idea is not “Rust catalogs are standard.” The portable idea is that a catalog state transition can be deterministic, auditable, replayable, and bound to engine facts without extra indirection.

The Turso-backed local store is also a LakeCat implementation choice, not an Iceberg term. Iceberg needs durable catalog state and atomic metadata-pointer movement; it does not name Turso, SQLite, PostgreSQL, FoundationDB, or any other storage engine. LakeCat uses the Rust turso crate for the local durable spine because it gives the project a fast embeddable store while keeping the store contract explicit. The important abstraction is `CatalogStore`: tables, views, namespaces, policies, storage profiles, idempotency records, pointer logs, audit rows, outbox rows, and replay markers must be read and written in ways that bind decoded data back to the requested tenant, warehouse, namespace, and table. Turso is the current local vehicle. The future portable concept is stronger than a database choice: exact catalog retry, atomic pointer movement, row-binding checks, transactional event identity, and replayable durable proof.

The Iceberg REST namespace and table paths are the compatibility floor. These are standard Iceberg parlance: catalog config, namespace, table identifier, metadata location, table metadata, requirements, optimistic commit, snapshot, manifest list, manifest, data file, delete file, partition

spec, sort order, and schema. A PySpark client should be able to use LakeCat as an Iceberg REST catalog without knowing what QGLake, TypeSec, Grust, or QueryGraph are. That is the line LakeCat must not cross. If normal table load or commit requires a non-standard QueryGraph endpoint, then LakeCat has stopped being an Iceberg-compatible catalog and has become a private application protocol. The LakeCat proof work should happen behind the route or beside the route, not in place of the route.

Commit CAS straddles standard language and LakeCat hardening. In Iceberg, a commit is optimistic: the client presents requirements that describe the table state it believes it is changing, and the catalog accepts the new metadata location only if the requirements still hold. LakeCat implements that as metadata-pointer compare-and-swap, which is standard in spirit: move from the previous metadata location to the new one only if the current state matches the expected state. LakeCat then adds proof around it: the idempotency record, metadata-object create-only posture, pointer log, audit row, outbox event, redacted conflict evidence, and replay admission verdict. CAS itself belongs to Iceberg-compatible commit correctness. The proof envelope is LakeCat's catalog reliability layer, and a stripped-down form of it could become a future optional catalog reliability profile.

Idempotency, pointer logs, audit/outbox, and replay validation are not the normal vocabulary a PySpark user brings to Iceberg. They are LakeCat's control-plane vocabulary. Idempotency means a retried commit with the same key must return the same accepted response or fail if the request drifted. Pointer logs mean operators and QueryGraph can reconstruct which metadata location was current after each accepted transition. Audit rows mean the catalog can explain who did what, under which principal and tenant, without exposing credentials or raw secrets. Outbox rows mean graph and lineage side effects are durable catalog facts that can be drained later rather than fragile request-path side effects. Replay validation means the drained event is checked again against a closed schema before Grust, OpenLineage, QGLake, or QueryGraph can inherit it. Those are extensions today. They are also among the best future proposal candidates because they can be described without adopting LakeCat's product names: "catalog event admission," "exact retry," "pointer history," and "transactional lineage outbox" are useful neutral ideas.

Governed scan and credential paths are LakeCat plus TypeSec plus Sail. Iceberg itself says an engine can use table metadata to plan a scan. It does not say that a catalog must return a TypeSec receipt, a purpose string, a policy hash, an allowed-column list, a mandatory predicate, a credential TTL cap, or a redacted proof of storage scope. LakeCat introduces those as additive governed surfaces because agents and untrusted principals should not default to broad object-store credentials. In the governed scan path, TypeSec-style policy logic decides the capability and restriction; LakeCat binds that receipt to the current pointer, request identity, tenant, and audit context; Sail turns the restriction into engine work. In the governed credential path, raw credential vending is intentionally narrower: it is an audited exception for a principal that is allowed to read directly, and its response must leave redacted prefix hashes, storage-profile hashes, TTL evidence, and secret-ref posture rather than leaking credential material. This is an extension today. The future proposal candidate is not "TypeSec in Iceberg"; it is a policy-engine-neutral governed access profile with proof-carrying scan and credential posture.

QueryGraph and QGLake handoff surfaces are product integration surfaces, not standard Iceberg semantics. QueryGraph needs bootstrap evidence, management inventory, table and view manifests, OpenLineage event hashes, graph import hashes, view receipt chains, credential posture, storage-profile posture, commit-history proof, replay verdicts, and artifact hashes. LakeCat should emit those anchors because QueryGraph needs a trustworthy substrate. But ordinary Iceberg clients should not need to produce or consume them. The standardization candidate is narrower than the QueryGraph feature set: event identity, lineage binding, replay-admissible payloads, view lifecycle proof, and commit proof could be phrased neutrally. Croissant, CDIF, OSI, ODRL, TypeDID workflow composition, QGLake artifact layout, Grust graph import shape, and agent workflow reasoning should remain QueryGraph, TypeSec, and Grust architecture unless and until a narrower neutral profile emerges.

Typed Iceberg v4 support is different again. If Iceberg v4 introduces or solidifies table-format semantics, those semantics are Iceberg work and engine work, not LakeCat product invention. LakeCat can preserve unknown JSON fields as a compatibility bridge, but passthrough is not interpretation. The durable direction is for Sail to own typed support: metadata structures, manifest and delete semantics, scan planning, metadata-as-data, commit validation, and view/table behavior. LakeCat can then advertise an honest state: standard REST compatibility now, JSON preservation where necessary, and typed Sail-backed proof as those engine APIs mature.

That is the extension answer. LakeCat has extensions in the product sense, because it extends what a catalog can prove. Most of those extensions should not be proposed as mandatory Iceberg behavior. The ones worth considering for Iceberg-adjacent discussion are the parts that can be stated without proper nouns: exact retry, pointer history, redacted conflict proof, transactional catalog events, replay-admissible lineage, governed scan proof, governed credential posture, and view lifecycle proof. The proper nouns are where the implementation is being proven. The neutral concepts are what might someday be shared.

12.2 Decision Matrix

	Concept	Iceberg meaning	LakeCat or Query-Graph meaning	Decision
Rust service/catalog spine		Iceberg does not specify an implementation language.	LakeCat keeps route handling, identity, tenancy, store transactions, Sail calls, TypeSec receipts, CAS, idempotency, audit, outbox, and replay close together in Rust.	Project implementation. The portable behavior is deterministic and replay-safe catalog state transition proof.
Turso-backed local store		Iceberg needs durable catalog state and atomic metadata-pointer movement, not a named database.	LakeCat uses the Rust turso crate as the durable local CatalogStore spine for tables, views, policies, storage profiles, idempotency, pointer logs, audit, outbox, and replay markers.	Project implementation. Generalize CAS, exact retry, pointer history, row binding, and transactional event identity, not Turso.
Iceberg REST namespace and table paths		Standard catalog compatibility: config, namespace operations, table load/create/drop, metadata locations, requirements, and optimistic commit.	LakeCat serves these paths while attaching server-side audit, outbox, and replay evidence.	Standard floor. Do not make normal clients depend on QueryGraph, TypeSec, Grust, QGLake, or LakeCat proof schemas.
Commit CAS		Optimistic metadata-pointer movement under commit requirements.	LakeCat binds the accepted transition to idempotency, pointer logs, audit, outbox, redacted conflict proof, and replay admission.	CAS is standard. The proof envelope is a future optional reliability-profile candidate.
Idempotency, pointer logs, audit/outbox, replay validation		Not generally standard table semantics, except insofar as they protect commit correctness.	LakeCat makes retries exact, pointer history inspectable, side effects recoverable, and down-stream provenance verifiable; centrally managed while all DDL changes happen through a single path, which is StatefulTableMetadata.	LakeCat extension today. Strong candidate for neutral catalog reliability and event-admission protection by default. Blending new, efficient practices, is sufficient to fully bring LakeCat into production life-cycle management and
QGLake client lineage/bootstrap/management/view/commit proof		Relational processing table semantics that enforce governance constraints. Typified by table receipts.	QGLake will ground-subject central authority while all DDL changes happen through a single path, which is StatefulTableMetadata.	Extension already follows. Blending new, efficient practices, is sufficient to fully bring LakeCat into production life-cycle management and

The standards answer is therefore not a blanket yes or no. LakeCat implements many extensions in the practical sense, because it extends what a catalog can prove. Only the parts that survive without LakeCat-specific names should be considered future Iceberg-adjacent proposals.

12.3 The Proper-Noun Test

The proper-noun test is the simplest review tool:

- If the concept requires the words LakeCat, QueryGraph, QGLake, TypeSec, Grust, Sail, or Turso to be meaningful, keep it as project architecture.
- If the concept can be described as behavior any Iceberg-compatible catalog or engine could implement, consider it a future optional-profile candidate.
- If the concept changes what a standard client must send or understand for an ordinary table load or commit, treat it as suspect until compatibility is proven.

That test keeps the proposal set small. Rust and Turso are excellent LakeCat choices, but they are not standards work. TypeSec receipts and QueryGraph import plans are essential product interfaces, but they should not become Iceberg requirements. Exact retry, pointer-history proof, redacted conflict proof, transactional event identity, replay-admissible evidence, governed scan proof, credential posture proof, view receipt chains, and lineage receipt binding are more plausible because they can be expressed without adopting the LakeCat stack.

12.4 Why The Engine Owns The Truth

The strongest technical reason to push work into Sail is that Iceberg table truth is not a set of easy strings. A governed proof that says “only these columns” is not reliable until it is mapped through Iceberg field ids, nested schema evolution, aliases, and projection rules. A proof that says “only these rows” is not reliable until the predicate has been interpreted by the same expression system that will plan the scan. A proof that says “these files” is not reliable until manifest metrics, residual predicates, partition transforms, delete files, sequence numbers, and snapshot context have been interpreted by the engine.

LakeCat can prove authority:

```
principal -> tenant -> warehouse -> namespace -> table
request -> TypeSec receipt -> current pointer -> CAS/idempotency
accepted state -> audit row -> outbox row -> replay admission
```

Sail should prove table work:

```
metadata -> schema field ids -> projection
metadata -> snapshots/manifests/statistics -> pruning
metadata -> equality/position deletes -> row visibility posture
metadata -> scan tasks / metadata-as-data / commit validation
```

When those two proofs are bound together, QueryGraph gets evidence that is both authorized and data-close. When LakeCat tries to do both jobs itself, it becomes a shadow engine: slower, less correct, and less reusable. Manifest metric decoding, delete planning, metadata-as-data, scan tasks, and typed v4 support should benefit Sail users, LakeCat governed reads, QueryGraph import, and the wider Rust lakehouse stack at the same time.

Sail is a particularly good engine boundary because it keeps the proof-heavy path Rust-to-Rust. LakeCat can receive a REST request, bind identity and tenancy, call TypeSec-style governance, ask Sail for typed table interpretation or a plan, commit through the Turso-backed CatalogStore, and persist replay proof without crossing a JVM sidecar, Python shim, or remote plugin boundary in the hot path. The public compatibility boundary remains Iceberg REST. The internal correctness boundary stays close to the Rust engine.

12.5 Workflow Consequences

For a PySpark user, the catalog remains ordinary. The user configures an Iceberg REST catalog, lists namespaces, loads a table, reads metadata, and commits with optimistic requirements. LakeCat's proof work stays server-side. The PySpark job does not need to understand TypeSec receipts, QGLake artifacts, Grust graph imports, or QueryGraph bootstrap.

For an operator, the catalog becomes inspectable. The operator can ask which principal moved a pointer, whether an idempotent retry drifted, which durable row was accepted, which outbox event is pending, which replay validator admitted the evidence, and whether downstream graph or OpenLineage projection came from accepted catalog state.

For a governed service, the catalog becomes enforceable. TypeSec narrows the request by purpose, capability, policy hash, allowed columns, mandatory predicate, TTL cap, and credential posture. LakeCat binds that decision to the current pointer and request identity. Sail turns the restriction into field-id projection, predicate binding, manifest/file pruning, delete posture, scan tasks, and plan evidence.

For an agent, the safest default becomes bounded work. Raw credentials are still possible for deliberately authorized principals, but they are audited exceptions. A restricted agent should usually receive Sail-planned tasks or a proof-backed fetch path whose scope, TTL, projection, predicate, and credential posture are already bound.

For QueryGraph, the catalog becomes a foundation rather than an application. LakeCat emits stable anchors: bootstrap hashes, management inventory hashes, policy-binding hashes, storage-profile hashes, commit-history hashes, view receipt-chain hashes, credential posture hashes, OpenLineage event hashes, graph/import hashes, artifact hashes, and replay verdicts. QueryGraph composes those anchors with Croissant, CDIF, OSI, ODRL, TypeDID, Grust graph state, and agent workflow state. LakeCat does not import QueryGraph. QueryGraph consumes LakeCat proof.

That separation is the release discipline. Standard clients get standard Iceberg behavior. Operators get durable Rust and Turso proof. Governed callers get TypeSec decisions bound to Sail engine work. QueryGraph gets a semantic handoff that is broad enough to build on and narrow enough not to fork Iceberg.

13 Ocelot: The 0.3 Release

Ocelot is the release where LakeCat's design becomes a public, reproducible claim. It does not contain the whole QueryGraph stack. It ships the compatible, durable, governed catalog foundation that the stack can trust, and it records the evidence needed to check that statement.

13.1 What shipped

The 0.3.0 catalog spine includes:

- standard Iceberg REST config, namespace, table-load, table-create, `listTables`, and table-commit paths, including warehouse-prefixed routing;
- the Rust service and `CatalogStore` seam with Turso-backed durable state and memory-store parity for embedded tests;
- metadata-pointer compare-and-swap, pointer history, idempotent replay, audit, and transactional outbox rows bound to the accepted catalog transition;
- TypeSec-gated scan and credential paths, Sail-planned work, Grust projection, OpenLineage drain, and `QueryGraph/QGLake` bootstrap evidence; and
- replay admission that rejects malformed, broadened, or secret-bearing durable evidence before downstream projection.

13.2 Stock-client conformance

Ocelot's compatibility claim is behavioral. An unmodified `PyIceberg RestCatalog` creates a namespace and table, appends real data, lists the table, loads it, and reads the rows back. Four release fixes made that ordinary round-trip possible: map-typed fields serialize as JSON objects; `listTables` is implemented; errors use real Iceberg `ErrorModel` exception types and statuses; and namespace creation/table creation follow stock-client semantics.

The commit path is honest in both feature postures. With `sail-local`, Sail applies supported Iceberg updates, including snapshot append. Without that engine integration, LakeCat rejects an update it cannot apply instead of returning success while leaving the table unchanged. Unknown or unsupported commit requirements also fail closed.

13.3 Governed reads and bounded credentials

The governed path is part of Ocelot because it is the reason LakeCat exists for agents. A restricted agent can ask for a governed read, receive a TypeSec receipt, get Sail-planned work instead of broad storage authority, and leave behind replayable scan, fetch, credential, management, view, and commit-history evidence. Trusted raw credential vending exists only as an audited exception with redacted storage-scope proof. Fetched scan replay treats the returned `plan-task` as evidence rather than arbitrary text: if it is present, it must be a non-empty LakeCat-issued token and it must not contain decorated location, query/fragment, or credential material before graph, OpenLineage, QGLake, or `QueryGraph` import can inherit the fetch proof. That keeps governed plan/fetch receipts about Sail-planned work, not a carrier for raw path or token claims. The same fetched replay path treats `stats-fields` as checked evidence when it is present: the array must be non-empty, duplicate-free, and bound to the effective stats fields before downstream proof can use it. The compact QGLake handoff proof now preserves that fetched side as its own `fetchRequestedStatsFields` and `fetchEffectiveStatsFields` evidence, so a handoff cannot prove only the planned stats narrowing while silently dropping what the fetch path actually returned. The local handoff script applies the same nonblank, duplicate-free array rule to projection, stats-field, and read-restriction allowed-column evidence before it writes that compact proof into the summary.

13.4 The handoff as release proof

The QueryGraph handoff is Ocelot’s acceptance proof, not a requirement for ordinary Iceberg clients. The QGLake workflow creates a bootstrap bundle, drains LakeCat lineage/outbox evidence, verifies replay, runs QueryGraph verification/import, and saves a handoff summary whose artifact hashes, table/view counts, standards, OpenLineage receipts, graph hashes, policy anchors, credential proof, view receipt chains, and commit history agree. If that handoff cannot be reproduced, LakeCat may still serve an Iceberg endpoint, but it has not reproduced the QueryGraph foundation claim.

The successful Ocelot gate ends with two plain statements:

QGLake handoff verified

LakeCat release-candidate checks passed from a clean tree.

The archived handoff is organized around four core files:

- lakecat-bootstrap.json, the governed catalog and semantic export;
- lineage-drain.json, the committed OpenLineage/outbox evidence;
- querygraph-import-plan.json, QueryGraph’s verified interpretation; and
- handoff-summary.json, the compact proof tying the other artifacts, captured command outputs, replay, graph projection, counts, and hashes together.

The saved handoff verifier output must also repeat the lineage-drain catalogConfigProof; omission is treated as proof failure, not as an unspecified optional field. Extra fields inside that repeated proof are also rejected, so an archived verifier sidecar cannot append a new endpoint or compatibility claim beside the validated catalog configuration evidence. Archived handoff file paths are part of that proof. LakeCat resolves artifact paths under the handoff summary directory before hashing or parsing them, and the verifier rejects relative traversal outside that bundle for both hash verification and captured-output semantic reads. A matching hash from another directory is not accepted as QGLake evidence. The artifact manifest itself is also closed: the primary artifacts object, nested capturedOutputs object, and individual bundle, lineage-drain, QueryGraph import-plan, and captured output artifact objects may carry only the path and SHA-256 evidence LakeCat checks. A saved summary cannot add an alternate hash, mirror artifact, or unverified capture beside an otherwise valid handoff bundle. The captured files themselves are closed at the root as well. Saved LakeCat replay output and QueryGraph verify/import output may carry only the fields LakeCat compares to the compact summary; a matching hash does not make an extra replay, QueryGraph, import, or application claim part of the proof. The same release gate treats raw view-lineage proof hashes as real digests, not placeholders: view replay receipts, tombstone view receipts, namespace receipt-chain hashes, and receipt-chain replay/OpenLineage hashes must be full SHA-256-shaped values before QGLake can archive them as accepted view proof.

13.5 Dependency provenance

Ocelot moves LakeCat onto the published **Grust 0.12.0**, **Lobster** and **TypeSec 0.12.0**, **Torcello** substrate. Grust’s grust-graph and grust-turso crates come from crates.io; TypeSec’s typesec facade comes from crates.io. Sail remains an explicit Cargo git dependency on the lakecat branch

of `github.com/querygraph/sail`, locked to the revision recorded in `Cargo.lock`. `QueryGraph 0.4 Sentinel` is an acceptance peer rather than a `LakeCat` build dependency; its importer consumes the `LakeCat 0.3` bundle contract and is exercised under `--locked`.

The dependency contract checks source as well as version. It rejects a quiet return to a local `Grust` path, a stale `TypeSec` line, a `Sail` source that no longer names the `lakecat` branch, or a `QueryGraph` importer that drops the shared `qglake-bundle` contract. The exact `Ocelot` table appears in *The Siblings and the Engine Path*.

`LakeCat`'s graph responsibility still stops at the catalog envelope. The sink rejects blank projection identity, non-object properties, and table-scoped labels without table identity before handing anything to `Grust`. Persistence, traversal, Cypher, and graph mutation semantics remain `Grust` work. Graph-store connect/bootstrap failures are reduced to `graph-store-path-hash` and `backend-error-hash` evidence so release logs do not capture raw database paths.

13.6 What Ocelot does not claim

Typed `Iceberg v4` semantics remain `Sail` work. `Ocelot` advertises the JSON passthrough bridge with `typed-sail=unavailable` rather than implying that a moving draft is fully modeled. Reusable graph taxonomy, traversal, algorithms, and stores remain in `Grust`. Capability composition, `Type-DID` envelopes, secure-agent proof, and richer policy semantics remain in `TypeSec`. `Croissant`, `CDIF`, `OSI`, `ODRL` application composition, and agent-facing reasoning remain in `QueryGraph`. `Cloud SDK` secret managers beyond the current provider seams are future credential backends.

Those are release boundaries, not omissions hidden behind an optimistic completion percentage. `Ocelot` proves the catalog foundation it ships and names the systems that own the rest.

13.7 Release and announcement posture

The announcement's proof-forward language comes from the same sources as the book: the release tag and changelog, the clean release-candidate record, the dependency lock and contract, and the `QGLake` handoff artifacts. Reader-facing docs and generated book formats are refreshed after the clean proof without pretending that a documentation rebuild is a new executable proof.

Tracked book artifacts are rebuilt with `docs/book/build.sh`. Release-candidate mode builds EPUB, PDF, MOBI, and HTML into a temporary `LAKECAT_B00K_DIST_DIR`, with temporary `Calibre` state, so format validation does not dirty the candidate commit or the operator's converter profile. Manual GitHub Actions remains a second environment; release-proof freshness is owned by the clean local gate.

13.8 Reproducing the proof

The release evidence is executable:

```
scripts/check-release-readiness.sh --release-candidate
scripts/qglake-handoff-local.sh
docs/book/build.sh
scripts/check-book-artifact-contract.sh docs/book/dist
scripts/check-local-dependency-contract.sh
```

The published Ocelot release-candidate proof ran on July 4, 2026 from clean head 6bfce1ef. It covered shell and dependency contracts, the default and all-feature Rust matrices, an out-of-tree book build, Grust-Turso projection, the QGLake handoff, QueryGraph locked verify/import, and the final clean-tree check. The immutable v0.3.0 tag records that release baseline; later documentation and artifact refreshes do not move it.

LakeCat also carries a smaller proof-freshness contract for the release docs themselves:

```
scripts/check-release-proof-contract.sh
```

That command verifies the active docs agree on the latest full release-candidate proof commit, that the proof commit is an ancestor of the current tree, and that any commits after it are limited to documentation and checked-in book artifact refresh. This avoids an infinite proof-update loop: the heavy gate proves the candidate, then a narrow documentation commit records the proof. If Rust code, manifests, workflows, dependency bridges, release scripts, or other executable behavior changes after the cited proof commit, the old proof is no longer enough and the full release-candidate gate must run again. The quick and ordinary full release-readiness gates surface that distinction without becoming release proof themselves. They print a non-failing freshness note when executable changes exist after the recorded proof commit, so a narrow development slice can still finish while the operator sees that final release proof requires a new `scripts/check-release-readiness.sh --release-candidate run`.

The contract requires a clean tree by default. While editing the proof contract or release docs, maintainers can run:

```
LAKECAT_RELEASE_PROOF_ALLOW_DIRTY=1 scripts/check-release-proof-contract.sh
```

That is only a local self-test mode. It still checks unstaged, staged, and untracked paths against the same post-proof allowlist, so a dirty Rust source file or manifest does not masquerade as documentation-only release evidence.

The full release-candidate gate runs the same contract in candidate mode:

```
LAKECAT_RELEASE_PROOF_CANDIDATE=1 scripts/check-release-proof-contract.sh
```

That mode is intentionally narrow. It still requires a clean tree and coherent active proof references, but it allows the current clean HEAD to become the next proof commit. After the heavy gate passes, the follow-up documentation and book artifact refresh records that new proof commit.

The already-published v0.1.0 tag is a baseline, not something to move. The already-published v0.2.0 tag is a baseline, not something to move. The already-published v0.2.1 tag is a baseline, not something to move. The already-published v0.3.0 tag is a baseline, not something to move. The workspace version is 0.3.0; post-Ocelot documentation and hardening stay under Unreleased until a deliberate future version bump. The release version contract checks that the workspace crates, release instructions, tracked book manifest, and versioned artifact links agree without turning a post-tag docs rebuild into a second Ocelot release.

The quick check is acceptable while landing a narrow slice:

```
scripts/check-release-readiness.sh --quick
```

The quick check is not a release claim, but it does validate the tracked docs/book/dist artifact contract so stale EPUB/PDF/MOBI/HTML deliverables are caught before a full release-candidate build regenerates temporary artifacts. The full gate is the release claim. It runs local dependency-contract checks, workflow-trigger checks, formatting, default workspace tests, feature matrix tests, Turso rows, Sail/TypeSec/Grust integration rows, all-features workspace tests, an out-of-tree book build with artifact validators, and the QGLake handoff proof. Cloud CI remains manual-only until that local proof is boring. A release should be cut from the state those commands prove, not from a hope that a remote runner will discover the truth later.

The handoff portion must stay locked in QueryGraph as well as in LakeCat. The local gate intentionally runs QueryGraph verification and import with Cargo's --locked mode so a stale QueryGraph dependency graph is a local failure, not a cloud surprise. When Grust or TypeSec publish new crates, the right sequence is to refresh the QueryGraph lockfile locally, prove lakecat-verify, lakecat-import, and QueryGraph tests under --locked, then rerun LakeCat's full release gate. That makes the QueryGraph import contract part of the release artifact instead of a best-effort downstream hope.

The current gate also proves a subtle QGLake contract that is easy to miss. The handoff readiness probe reads GET /catalog/v1/config, and that read is a catalog event. It must therefore carry the same QGLake agent identity as the rest of the handoff, otherwise a harmless-looking readiness curl becomes an anonymous catalog.config-read proof in the drained evidence. The handoff script binds that config-read proof to the agent, writes a canonical handoff-verifier output artifact, hashes it into the summary before strict verification, and keeps management list proof, policy-upsert proof, and storage-profile upsert proof in the distinct sections the verifier compares. That is the local substitute for cloud hope: the artifact graph is checked before the summary is accepted.

For Ocelot, replay admission is part of the local proof rather than a best-effort projection filter. The service closes event payloads before graph, OpenLineage, QGLake, or QueryGraph import can inherit them; the compact handoff verifier then repeats the same checks against saved artifacts. Appendix A lists the replay-admission contract, Appendix B captures credential redaction, Appendix C captures view receipt chains, and Appendix D maps the QGLake handoff proof.

14 Appendix A: Replay Admission Matrix

Replay admission is the line between durable catalog fact and downstream proof. An outbox row may be stored, but it is not allowed to become graph, OpenLineage, QGLake, or QueryGraph evidence until LakeCat has checked its shape, scope, authority, and redaction posture. The matrix below keeps the best review version of those checks in one place.

Event family	Must prove	Rejects	Why it matters
catalog.config-read	Warehouse scope, v4 bridge defaults, overrides, advertised standard REST endpoints, governed plan/fetch/credential endpoints, bootstrap and lineage-drain endpoints, authorization receipt	Duplicate config keys, unsupported lakecat.format.v4* claims, missing standard or governed endpoints, extra unverified config fields	QueryGraph imports the same compatibility contract LakeCat advertised, without turning config into a free-form claims bag
Namespace lifecycle	Warehouse, namespace path or components, matching lifecycle action, valid principal	Malformed namespace paths, count drift, action drift, unclosed wrapper fields	Standard catalog discovery stays attributable and replayable
Table lifecycle	Root table identity, warehouse/namespace/table hints matching that identity, action receipt, soft-delete format evidence when present	Scope drift, action drift, conflicting format-version/format_version aliases, undecorated-location violations	Table proof follows the Iceberg object it claims to describe
Table commit	Table scope, authorization receipt, pointer transition, principal, idempotency/request/response/policy hashes, format, snapshot, timestamp	Credential-bearing or decorated metadata-location evidence, unclosed nested commit fields, forged event type	Commit proof stays tied to the accepted pointer movement without leaking object paths
Governed scan and fetch	Principal, table-plan-scan action, requested/effective projections, requested/effective stats fields, required filters, matching top-level and receipt read restrictions, non-empty full policy hashes, past is more restrictive	Actorless scan proof, action drift, missing or widening projection/filter evidence, placeholder policy hashes, non-LakeCat or credential-bearing plan-task values	A saved scan proof shows the same TypeSec restriction Sail planned and LakeCat replayed
Warehouse plan/fetch/commit	Warehouse plan/fetch/commit action, table-plan-scan action, requested/effective projections, requested/effective stats fields, required filters, matching top-level and receipt read restrictions, non-empty full policy hashes, past is more restrictive	Malformed plan/fetch/commit fields, malformed table-plan-scan fields, malformed requested/effective projections, malformed requested/effective stats fields, malformed required filters, malformed matching top-level and receipt read restrictions, malformed non-empty full policy hashes, malformed past is more restrictive	Warehouse plan/fetch/commit action, table-plan-scan action, requested/effective projections, requested/effective stats fields, required filters, matching top-level and receipt read restrictions, non-empty full policy hashes, past is more restrictive

The pattern is deliberately repetitive in code but not in the narrative: every event must be scoped, actor-bound, action-bound, closed over the fields LakeCat actually verifies, and redacted before it becomes downstream evidence.

15 Appendix B: Credential And Redaction Contract

Credential handling is where catalog convenience can become security risk. LakeCat's posture is therefore narrow: governed Sail-planned reads are the default for agents, while raw credentials are explicit audited exceptions.

Surface	Accepted shape	Redaction rule	Failure posture
Warehouse storage root	Plain URI rooted in the warehouse; no query, fragment, userinfo, or dot traversal	Return warehouse-storage-root-hash rather than raw root on errors	Reject before memory or Turso persistence
Storage-profile prefix	Provider-compatible URI prefix, longest-prefix match, no ambiguous tie	Return storage-profile-prefix-hash or location-prefix-hash	Fail closed rather than guessing a credential root
Public config	Non-secret string hints such as region, endpoint labels, or purpose	Return public-config-key-hash for rejected keys/values	Reject secret-looking keys, values, and LakeCat-reserved proof keys
Secret reference	Clean external locator such as vault://, aws-sm://, gcp-sm://, or azure-kv://	Store and replay secret-ref-present, provider, and full secret-ref-hash	Reject decorated URIs, unsupported schemes, traversal, userinfo, query, fragment
Resolver dispatch	TypeSec authorizes the exact secret-ref resource before lookup	Resolver failures report provider plus secret/error hashes	Denied decisions do not call the resolver
Returned credentials	Prefix remains inside the selected profile; TTL honors policy cap; LakeCat-owned evidence is canonical	Replay prefix hashes and issuer-config hashes, never raw secret material	Reject widened prefixes before recording credentials.vend-attempted
Restricted agent branch	Zero credentials, explicit block reason, raw exception false, TTL cap present	Empty credential-response evidence is explicit proof, not absence	Agent is steered to governed Sail-planned reads
Trusted human branch	Positive credential count only when policy allows raw exception	Returned prefixes are full hashes; exception reason is recorded	Raw credentials remain auditable exceptions, not the default path

The important distinction is that storage-profile management proves a credential root exists, while credential vending proves whether a caller received raw storage reach. QueryGraph can

reason over both without seeing the raw root, secret URI, backend error, token, or returned credential material.

16 Appendix C: View Receipt-Chain Rules

Ocelot does not claim full Iceberg view-history semantics. It keeps a compact, catalog-owned receipt chain so QueryGraph can prove which view version was exported, updated, dropped, or tombstoned. Richer view semantics remain a Sail-aligned future target.

Rule	Required evidence	Rejects
Stable identity	lakecat:view:<warehouse>:<namespace> plus matching component fields	Stable IDs that do not match warehouse/namespace/name
First receipt	Version 1 upsert, no previous version, no previous receipt hash	Zero-version chains, first tombstones, forged first previous-link fields
Later upsert	Previous receipt hash points at the prior receipt; version advances by exactly one	Skipped versions, unsupported operations, broken previous links
Drop tombstone	Operation drop, previous receipt hash set, view-version remains the deleted version	Drops that advance the version or omit the previous receipt
Recreate	New upsert links after the tombstone and advances to the next version	A second unrelated version-1 chain for the same stable view id
Chain header	receipt-count, latest version, latest operation, and tombstone flag agree with the last receipt	Headers that shed receipts or present a forged head
Hash coverage	receipt-hashes, drop-receipt-hashes, and chain-hashes exactly cover the nested receipt bodies and chains	Convenient full-looking hashes that do not match the embedded chain
Durable row binding	Turso row warehouse/name-space/view columns agree with decoded receipt JSON	A valid-looking receipt body riding a corrupted row index into another view
QueryGraph handoff	Accepted view receipt hashes match bootstrap view artifacts and compact handoff proof	Spliced receipt arrays from another view or another run

The receipt chain is intentionally compact: it proves catalog state changes and tombstones, not full SQL engine view semantics. That is enough for QueryGraph to reject a bootstrap that lost or borrowed view proof while LakeCat keeps richer view interpretation out of the catalog.

17 Appendix D: QGLake Handoff Proof Map

The QGLake handoff is a compact acceptance artifact over several raw artifacts: the LakeCat bootstrap bundle, lineage-drain response, QueryGraph import plan, captured verifier output, captured replay output, and service log. The summary is useful only because every compact field can be traced back to a parsed raw artifact and a replay-admitted catalog event.

Proof section	Binds	Compared against
requestIdentityProof	Principal subject/kind, identity source/state, lineage-drain receipt, TypeDID hash slots	Captured LakeCat replay request identity and saved handoff-verifier output
queryGraphBootstrapProof	Bundle hash, graph hash, OpenLineage hash, QueryGraph import hash, standards, artifact counts, policy count, view receipt hashes	Bootstrap replay event, bundle manifest, QueryGraph verify/import output
governedScanProof	Planned/fetched task counts, projections, stats fields, filters, read restrictions, policy hashes, TTL cap	Scan replay sections and operator replay text
tableCommitHistoryProof	Commit count, positive increasing sequence numbers, duplicate-free commit hashes, replay/OpenLineage hashes	Pointer-log read replay and compact summary
viewReceiptChainProof	Accepted view versions, receipt hashes, receipt-chain hashes, tombstone receipts, verified-chain counts	View replay sections, bootstrap view artifacts, receipt-chain reads
managementProof	Server/project/warehouse/policy/storage-profile counts, ID arrays, graph event counts, replay/OpenLineage hashes	Management-list replay sections and QueryGraph bootstrap policy count
policyUpsertProof	Policy id, ODRL hash, principal/action, graph events, replay/OpenLineage hashes	Policy upsert replay and policy-list proof
storageProfileUpsertProof	Profile id, provider, issuance mode, location-prefix hash, secret-ref posture, principal/action	Storage-profile upsert replay and credential branches
credentialVendingProof	Restricted-agent block, trusted-human exception, credential counts, prefix hashes, TTL caps, storage-profile anchor	Credential replay sections and operator replay text
Artifact manifest	Bundle, lineage drain, import plan, captured outputs, service log, saved ⁵⁶ handoff verifier output	Recomputed file hashes under the handoff summary directory

The verifier enforces three global rules. First, artifact paths resolve under the handoff summary directory before hashing or parsing, so a summary cannot splice matching bytes from outside the bundle. Second, all integrity anchors are full sha256:-prefixed 64-hex digests, not readable placeholders. Third, saved sidecars are closed over the schema LakeCat compares; a hash match does not make extra unverified fields part of the proof.

This is why the handoff is a release acceptance artifact instead of a convenient log bundle. It proves that LakeCat replay, graph projection, OpenLineage, QueryGraph verification, and QueryGraph import all agreed on the same catalog state.

18 Future Work

LakeCat is a direction more than a single release. The next stage should keep the architecture more true without pushing engine, graph, or security semantics back into the catalog.

Marciana v1 is complete as a durable local-service proof; implementing agent memory is therefore not a LakeCat backlog item. LakeCat's future opportunity is narrower: let a downstream QueryGraph workflow bind an allowed answer to the catalog receipt, scan-plan, snapshot, bootstrap, and OpenLineage hashes that justify it, so those LakeCat-owned anchors become Marciana provenance without copying raw table data into the catalog or making LakeCat responsible for recall.

1. Point the Sail git dependency at upstream (or a published crate) only after the required helpers land there, then retire the querygraph/sail lakecat branch and its dependency-contract assertions.
2. Keep pushing reusable table-status, planning, metadata-as-data, manifest, delete, and commit helpers into Sail.
3. Treat typed Iceberg v4 as a Sail project after formal specification adoption. The active v4 draft includes relative metadata locations, but LakeCat should not turn a moving draft into catalog behavior; its bridge remains explicit and JSON-only until Sail can model and plan the adopted semantics.
4. Make the transactional outbox the only path for graph and lineage side effects, while keeping graph persistence, traversal, and Cypher in Grust.
5. Expand TypeSec-backed capability checks until every privileged path is unbypassable, with raw credentials remaining an explicit audited exception.
6. Keep OSI, Croissant, CDIF, ODRL composition, and agent workflows in QueryGraph. Keep QueryGraph 0.4 Sentinel's importer aligned with LakeCat qglake-bundle/lakecat-core 0.3.0 and the Grust/TypeSec 0.12.0 substrate while preserving the locked receipt-chain and graph-handoff proof.
7. Prove the bootstrap bundle through QueryGraph import on every meaningful public-surface change, and keep local release evidence ahead of cloud CI.
8. Define an optional handoff profile that binds governed answer text to LakeCat receipt and lineage hashes before qg-rust writes it to Marciana, while keeping capability minting, content rehydration, and recall wholly in TypeSec and QueryGraph.

That optional binding must preserve the upstream v1 limits rather than turning them into LakeCat claims. `MemoryId::next()` is process-local, so a new write after `qg-rust` restarts can collide with an existing persisted id; Grust pushes down space equality but evaluates the other memory query dimensions through the shared matcher; the reference vector index is in-process and nonpersistent; and there is no hosted multi-tenant control plane yet. Those are upstream Marciana scale contracts to resolve before a production LakeCat-to-memory profile can claim restart-safe writes, full query pushdown, durable semantic ranking, or hosted tenancy.

The payoff is a catalog foundation that feels ordinary to Iceberg clients and rich to QueryGraph. Tables remain portable. Policies become enforceable. Graph and lineage become replayable. Sail plans close to the data. QueryGraph gets a trustworthy substrate for the next version of its lakehouse intelligence.